

FINETUNED LANGUAGE MODELS ARE ZERO-SHOT LEARNERS

Jason Wei* Maarten Bosma* Vincent Y. Zhao* Kelvin Guu* Adams Wei Yu
Brian Lester Nan Du Andrew M. Dai Quoc V. Le

Google Research

ABSTRACT

This paper explores a simple method for improving the zero-shot learning abilities of language models. We show that *instruction tuning*—finetuning language models on a collection of datasets described via instructions—substantially boosts zero-shot performance on unseen tasks.

We take a 137B parameter pretrained language model and instruction tune it on over 60 NLP datasets verbalized via natural language instruction templates. We evaluate this instruction-tuned model, which we call FLAN, on unseen task types. FLAN substantially improves the performance of its unmodified counterpart and surpasses zero-shot GPT-3 on 20 of 25 datasets that we evaluate. FLAN even outperforms few-shot GPT-3 by a large margin on ANLI, RTE, BoolQ, AI2-ARC, OpenbookQA, and StoryCloze. Ablation studies reveal that number of datasets and model scale are key components to the success of instruction tuning.

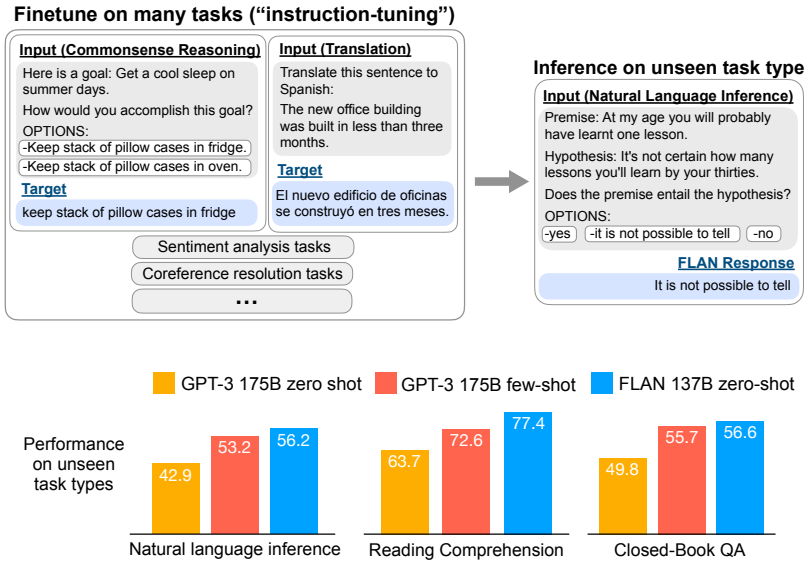


Figure 1: Top: overview of instruction tuning and FLAN. Instruction tuning finetunes a pretrained language model on a mixture of tasks phrased as instructions. At inference time, we evaluate on an unseen task type; for instance, we could evaluate the model on natural language inference (NLI) when no NLI tasks were seen during instruction tuning. Bottom: performance of zero-shot FLAN on unseen task types, compared with zero-shot GPT-3 and few-shot. NLI datasets: ANLI R1–R3, CB, RTE. Reading comprehension datasets: BoolQ, MultiRC, OBQA. Closed-book QA datasets: ARC-easy, ARC-challenge, NQ, TriviaQA.

*Lead contributors. Author contributions listed at end of paper.

1 INTRODUCTION

Language models (LMs) at scale, such as GPT-3 (Brown et al., 2020), have been shown to perform few-shot learning remarkably well. They are less successful at zero-shot learning, however. For example, GPT-3’s zero-shot performance is much worse than few-shot performance on tasks such as reading comprehension, question answering, and natural language inference. One potential reason is that, without few-shot exemplars, it is harder for models to perform well on prompts that are not similar to the format of the pretraining data.

In this paper, we explore a simple method to improve the zero-shot performance of large language models, which would expand their reach to a broader audience. We leverage the intuition that NLP tasks can be described via natural language instructions, such as “*Is the sentiment of this movie review positive or negative?*” or “*Translate ‘how are you’ into Chinese.*” We take a pretrained language model of 137B parameters and perform *instruction tuning*—finetuning the model on a mixture of more than 60 NLP datasets expressed via natural language instructions. We refer to this resulting model FLAN, for Finetuned LLanguage Net.

To evaluate the zero-shot performance of FLAN on unseen tasks, we group NLP datasets into clusters based on their task types and hold out each cluster for evaluation while instruction tuning FLAN on all other clusters. For example, as shown in Figure 1, to evaluate FLAN’s ability to perform natural language inference, we instruction tune the model on a range of other NLP tasks including commonsense reasoning, translation, and sentiment analysis. As this setup ensures that FLAN has not seen any natural language inference tasks in instruction tuning, we then evaluate its ability to perform zero-shot natural language inference.

Our evaluations show that FLAN substantially improves the zero-shot performance of the base 137B-parameter model. FLAN’s zero-shot also outperforms 175B-parameter GPT-3’s zero-shot on 20 of 25 datasets that we evaluate, and even outperforms GPT-3’s few-shot by a large margin on ANLI, RTE, BoolQ, AI2-ARC, OpenbookQA, and StoryCloze. In ablation studies, we find that increasing the number of task clusters in instruction tuning improves performance on unseen tasks and that the benefits of instruction tuning emerge only with sufficient model scale.

Instruction tuning is simple method that, as depicted in Figure 2, combines appealing aspects of both the pretrain–finetune and prompting paradigms by using supervision via finetuning to improve language model’s responses to inference-time text interactions. Our empirical results demonstrate promising abilities of language models to perform tasks described purely via instructions.

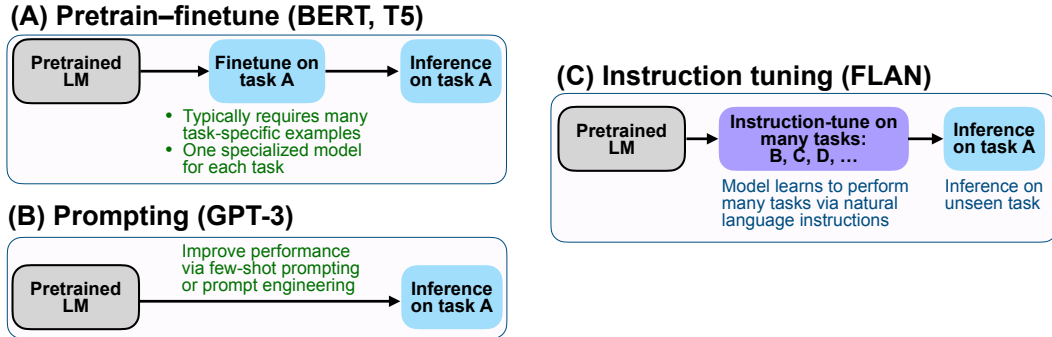


Figure 2: Comparing instruction tuning with pretrain–finetune and prompting.

2 FLAN: INSTRUCTION TUNING IMPROVES ZERO-SHOT LEARNING

The motivation of instruction tuning is to improve the ability of language models to respond to NLP instructions. The idea is that by using supervision to teach an LM to perform tasks described via instructions, it will learn to follow instructions and do so even for unseen tasks. To evaluate the model’s performance on unseen tasks, we group datasets into clusters by task type and hold out each task cluster for evaluation while instruction tuning on all remaining clusters.

2.1 TASKS & TEMPLATES

As creating an instruction tuning dataset with many tasks from scratch would be resource-intensive, we transform existing datasets from the research community into an instructional format. We aggregate 62 text datasets that are publicly available on Tensorflow Datasets, including both language understanding and language generation tasks, into a single mixture. Figure 3 shows these datasets; each dataset is categorized into one of twelve task clusters, for which datasets in a given cluster are of the same task type. Descriptions, sizes, and examples of each dataset are shown in Appendix F.

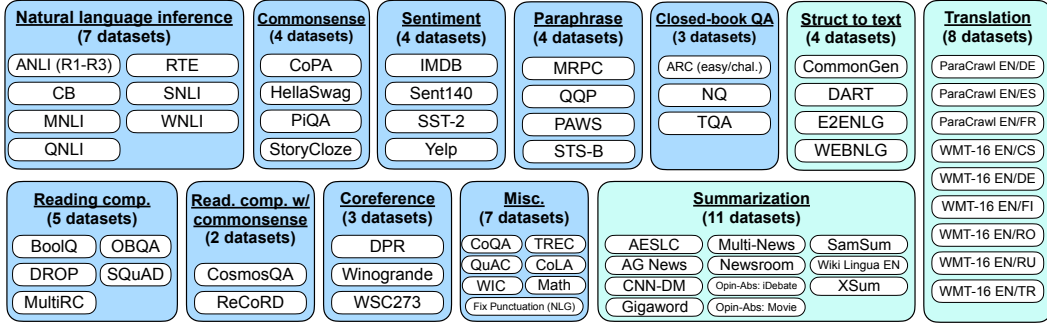


Figure 3: Datasets and task clusters used in this paper (NLU tasks in blue; NLG tasks in teal).

For each dataset, we manually compose ten unique templates that use natural language instructions to describe the task for that dataset. While most of the ten templates describe the original task, to increase diversity, for each dataset we also include up to three templates that “turned the task around,” (e.g., for sentiment classification we include templates asking to generate a movie review). We then instruction tune a pretrained language model on the mixture of all datasets, with examples in each dataset formatted via a randomly selected instruction template for that dataset. Figure 4 shows multiple instruction templates for a natural language inference dataset.

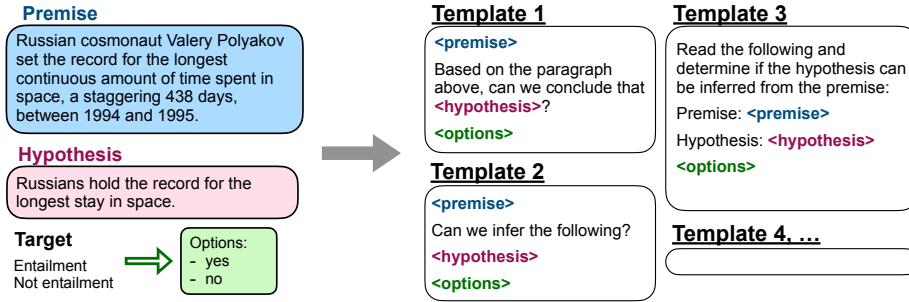


Figure 4: Multiple instruction templates describing a natural language inference task.

2.2 EVALUATION SPLITS

We are interested in how FLAN performs on tasks not seen in instruction tuning, and so it is crucial to define what counts as an unseen task. Whereas some prior work defines unseen tasks by disallowing the same dataset to appear in training, we use a more conservative definition that leverages the task clusters from Figure 3. In this work, we only consider dataset \mathcal{D} unseen at evaluation time if no datasets from any task clusters that \mathcal{D} belongs to were seen during instruction tuning. For instance, if \mathcal{D} is an entailment task, then no entailment datasets appeared in instruction tuning, and we instruction-tuned on all other clusters.¹ Hence, to evaluate zero-shot FLAN on c task clusters, we instruction tune c models, where each model holds out a different task cluster for evaluation.

¹When evaluating on the read. comp. with commonsense cluster, both read. comp. and commonsense reasoning were dropped from instruction tuning. Conversely, the read. comp. with commonsense cluster was not used for instruction tuning when evaluating on read. comp. or commonsense reasoning. We also drop the paraphrase cluster from instruction tuning when evaluating on NLI tasks and vice-versa.

2.3 CLASSIFICATION WITH OPTIONS

The output space for a given task is either one of several classes (classification) or free text (generation). As FLAN is an instruction-tuned version of a decoder-only language model, it naturally responds in free text, and so no further modifications are needed for generation tasks.

For classification tasks, prior work (Brown et al., 2020) has used a *rank classification* approach where, for example, only two outputs (“yes” and “no”) are considered and the higher probability one is taken as the model’s prediction. Though this procedure is logically sound, it is imperfect in that the probability mass for answers may have an undesired distribution among ways of saying each answer (e.g., a large number of alternative ways of saying “yes” may lower the probability mass assigned to “yes”). Therefore, we include *an options suffix*, in which we append the token `OPTIONS` to the end of a classification task along with a list of the output classes for that task. This makes the model aware of which choices are desired when responding to classification tasks. Example use of options is shown in the NLI and commonsense examples in Figure 1.

2.4 TRAINING DETAILS

Model architecture and pretraining. In our experiments, we use a *dense left-to-right, decoder-only transformer language model of 137B parameters*. This model is pretrained on a collection of web documents (including those with computer code), dialog data, and Wikipedia, tokenized into 2.81T BPE tokens with a 32k vocabulary using the SentencePiece library (Kudo & Richardson, 2018). Around 10% of the pretraining data was non-English. This dataset is not as clean as the GPT-3 training set and also *has a mixture of dialog and code*, and so we expect the zero and few-shot performance of this pretrained LM on NLP tasks to be slightly lower. We refer to this pretrained model as *Base LM*. Base LM was also previously used for program synthesis (Austin et al., 2021).

Instruction tuning procedure. FLAN is the instruction-tuned version of Base LM. Our instruction tuning pipeline *mixes all datasets and randomly samples from each dataset*. To balance the different sizes of datasets, we limit the number of training examples per dataset to 30k and follow the *examples-proportional mixing scheme* (Raffel et al., 2020) with a mixing rate maximum of 3k. We finetune all models for 30k gradient steps with a batch size of 8,192 tokens using the *Adafactor Optimizer* (Shazeer & Stern, 2018) with a learning rate of $3e-5$. The input and target sequence lengths used in finetuning are 1024 and 256, respectively. We use *packing* (Raffel et al., 2020) to combine multiple *training examples into a single sequence, separating inputs from targets using a special EOS token*. For all evaluations, we report results on the final checkpoint trained for 30k steps.

3 RESULTS

We evaluate FLAN on natural language inference, reading comprehension, closed-book QA, translation, commonsense reasoning, coreference resolution, and struct-to-text. As described in §2.2, we evaluate on unseen tasks by grouping datasets into task clusters and holding out each cluster for evaluation while instruction tuning on all remaining clusters (i.e., each evaluation task cluster uses a different checkpoint). For each dataset, we report the mean and standard deviation of performance on all templates, which proxies the expected performance given a typical natural language instruction. As a dev set is sometimes available for manual prompt engineering (Brown et al., 2020), for each dataset we also report the test set performance using the template with the best dev set performance.

For comparison, we show the zero-shot and few-shot performance of GPT-3 175B, as reported in their paper. With the best dev template, zero-shot FLAN outperforms zero-shot GPT-3 on 20 out of 25 tasks and even surpasses GPT-3’s few-shot performance on 10 tasks. We also report zero and few-shot results from Base LM using the same prompts as GPT-3 (as Base LM is not suitable for natural instructions without instruction tuning). This baseline provides the most direct ablation of how much instruction tuning helps. Instruction tuning significantly improves Base LM on most tasks. In terms of supervised models, we report the performance of T5-11B (Raffel et al., 2020) when available; otherwise, we show the performance of BERT-large (Devlin et al., 2019).

Overall, we observe that *instruction tuning is very effective on tasks naturally verbalized as instructions* (e.g., NLI, QA, translation, struct-to-text) and is less effective on tasks directly formulated as language modeling, where instructions would be largely redundant (e.g., commonsense reasoning and coreference resolution tasks that are formatted as finishing an incomplete sentence or paragraph).

3.1 NATURAL LANGUAGE INFERENCE

Table 1 shows results for natural language inference (NLI), where a model must determine whether a hypothesis is true given some premise. FLAN has strong performance in all cases. Despite high variance among different templates for CB (De Marneffe et al., 2019) and RTE (Wang et al., 2018), FLAN outperforms zero and few-shot GPT-3 for four datasets by a large margin without any prompt engineering. With the best dev template, FLAN outperforms few-shot GPT-3 on all five datasets.

As noted by Brown et al. (2020), perhaps one reason why GPT-3 struggles with NLI is that NLI examples are unlikely to have appeared naturally in an unsupervised training set and are thus awkwardly phrased as a continuation of a sentence. For FLAN, we phrase NLI as the more natural question “Does <premise> mean that <hypothesis>?” and achieve much higher performance.

	NATURAL LANGUAGE INFERENCE				
	ANLI-R1 acc.	ANLI-R2 acc.	ANLI-R3 acc.	CB acc.	RTE acc.
Supervised model	57.4 ^b	48.3 ^b	43.5 ^b	96.8 ^a	92.5 ^a
Base LM 137B zero-shot	39.6	39.9	39.3	42.9	73.3
· few-shot	39.0	37.5	40.7	34.8	70.8
GPT-3 175B zero-shot	34.6	35.4	34.5	46.4	58.9
· few-shot	36.8	34.0	40.2	82.1	70.4
FLAN 137B zero-shot					
- average template	47.7 ^{▲10.9} std=1.4	43.9 ^{▲8.5} std=1.3	47.0 ^{▲6.8} std=1.4	64.1 ^{↑17.7} std=14.7	78.3 ^{▲7.9} std=7.9
- best dev template	46.4 ^{▲9.6}	44.4 ^{▲9.0}	48.5 ^{▲8.3}	83.9 ^{▲1.8}	84.1 ^{▲13.9}

Table 1: Results on NLI. For FLAN, we report both the average of up to ten templates, as well as the best dev template. The triangle ▲ indicates improvement over few-shot GPT-3. The up-arrow ↑ indicates improvement only over zero-shot GPT-3. ^aT5-11B, ^bBERT-large.

3.2 READING COMPREHENSION & CLOSED-BOOK QA

Results for reading comprehension, where models are asked to answer a question about a provided passage, are shown in Table 2. FLAN outperforms GPT-3 by a large margin on BoolQ (Clark et al., 2019a) and OBQA (Mihaylov et al., 2018), although Base LM already achieves high performance on BoolQ. On MultiRC (Khashabi et al., 2018), the performance of FLAN when using the best dev template is slightly higher than few-shot GPT-3.

	READING COMPREHENSION			CLOSED-BOOK QA			
	BoolQ acc.	MultiRC F1	OBQA acc.	ARC-e acc.	ARC-c acc.	NQ EM	TQA EM
Supervised model	91.2 ^a	88.2 ^a	85.4 ^a	92.6 ^a	81.1 ^a	36.6 ^a	60.5 ^a
Base LM 137B zero-shot	81.0	60.0	41.8	76.4	42.0	3.2	21.9
· few-shot	79.7	59.6	50.6	80.9	49.4	22.1	63.3
GPT-3 175B zero-shot	60.5	72.9	57.6	68.8	51.4	14.6	64.3
· few-shot	77.5	74.8	65.4	70.1	51.5	29.9	71.2
FLAN 137B zero-shot							
- average template	80.2 ^{▲2.7} std=3.1	74.5 ^{↑2.4} std=3.7	77.4 ^{▲12.0} std=1.3	79.5 ^{▲8.6} std=0.8	61.7 ^{▲10.2} std=1.4	18.6 ^{↑4.0} std=2.7	66.5 ^{↑2.2} std=2.6
- best dev template	82.9 ^{▲5.4}	77.5 ^{▲2.7}	78.4 ^{▲13.0}	79.6 ^{▲8.7}	63.1 ^{▲11.6}	20.7 ^{↑6.1}	68.1 ^{↑3.8}

Table 2: Results on reading comprehension and closed-book question answering. For FLAN, we report both the average of up to ten templates, as well as the best dev template. The triangle ▲ indicates improvement over few-shot GPT-3. The up-arrow ↑ indicates improvement only over zero-shot GPT-3. ^aT5-11B.

Table 2 also shows results for closed-book QA, which asks models to answer questions about the world without access to specific information containing the answer. On ARC-easy and ARC-challenge

(Clark et al., 2018), FLAN substantially outperforms both zero-shot and few-shot GPT-3. For Natural Questions (NQ; Lee et al., 2019; Kwiatkowski et al., 2019) and TQA (Joshi et al., 2017), FLAN outperforms zero-shot GPT-3 but does not outperform few-shot GPT-3.

3.3 TRANSLATION

Similar to GPT-3, the training data for Base LM is around 90% English and includes some text in other languages that was not specifically used to train the model to perform machine translation. We also evaluate FLAN’s performance on machine translation for the three datasets evaluated in the GPT-3 paper: French–English from WMT’14 (Bojar et al., 2014), and German–English and Romanian–English from WMT’16 (Bojar et al., 2016).

These results are shown in Table 3. For Base LM, zero-shot translation had low performance, but few-shot translation results were comparable with GPT-3; FLAN outperforms few-shot Base LM in five of the six evaluations. Compared with GPT-3, FLAN outperforms zero-shot GPT-3 for all six evaluations, though it underperforms few-shot GPT-3 in most cases. Similar to GPT-3, FLAN shows strong results for translating into English and compares favorably against supervised translation baselines. Translating from English into other languages, however, was relatively weaker, as might be expected given that FLAN uses an English sentencepiece tokenizer and that the majority of pretraining data is English.

	TRANSLATION					
	French		German		Romanian	
	En→Fr BLEU	Fr→En BLEU	En→De BLEU	De→En BLEU	En→Ro BLEU	Ro→En BLEU
Supervised model	45.6 ^c	35.0 ^d	41.2 ^e	38.6 ^f	38.5 ^g	39.9 ^g
Base LM 137B zero-shot	11.2	7.2	7.7	20.8	3.5	9.7
· few-shot	31.5	34.7	26.7	36.8	22.9	37.5
GPT-3 175B zero-shot	25.2	21.2	24.6	27.2	14.1	19.9
· few-shot	32.6	39.2	29.7	40.6	21.0	39.5
FLAN 137B zero-shot						
- average template	32.0 \uparrow 6.8 std=2.0	35.6 \uparrow 14.4 std=1.5	24.2 std=2.7	39.4 \uparrow 12.2 std=0.6	16.9 \uparrow 2.8 std=1.4	36.1 \uparrow 16.2 std=1.0
- best dev template	34.0 \blacktriangle 1.4	36.5 \uparrow 15.3	27.0 \uparrow 2.4	39.8 \uparrow 12.6	18.4 \uparrow 4.3	36.7 \uparrow 16.7

Table 3: Translation results (BLEU) for WMT’14 En/Fr and WMT’16 En/De and En/Ro. For FLAN, we report both the average of up to ten templates, as well as the best dev template. ^cEduonov et al. (2018), ^dDurrani et al. (2014), ^eWang et al. (2019b), ^fSennrich et al. (2016), ^gLiu et al. (2020). The triangle \blacktriangle indicates improvement over few-shot GPT-3. The up-arrow \uparrow indicates improvement only over zero-shot GPT-3.

3.4 ADDITIONAL TASKS

Although we see strong results for the above task clusters, one limitation with instruction tuning is that it does not improve performance for many language modeling tasks (e.g., commonsense reasoning or coreference resolution tasks that are formulated as sentence completions). To summarize the full result shown in Table 4 in Appendix A, for seven commonsense reasoning and coreference resolution tasks, FLAN only outperforms Base LM on three of the seven tasks, and, with the best dev template, only outperforms zero-shot GPT-3 in two of the seven tasks. This negative result indicates that when the downstream task is the same as the original language modeling pre-training objective (i.e., in cases where instructions are largely redundant), instruction tuning is not useful.

Finally, we report results for the sentiment analysis, paraphrase detection, and struct to text task clusters, as well as other datasets for which GPT-3 results are not available, in Table 5 and Table 6 in Appendix A. Generally, we see that **zero-shot FLAN substantially outperforms zero-shot Base LM** and is comparable with or better than few-shot Base LM, though FLAN’s performance is still far from that of supervised models on most tasks.

4 ABLATION STUDIES & FURTHER ANALYSIS

4.1 NUMBER OF INSTRUCTION TUNING CLUSTERS

As the core question of our paper asks how instruction tuning improves a model’s zero-shot performance on unseen tasks, in this first ablation we examine how performance is affected by the number of clusters and tasks used in instruction tuning. For this setup, we hold out NLI, closed-book QA, and commonsense reasoning as evaluation clusters, and use the seven remaining clusters for instruction tuning.² We show results for 1-7 instruction tuning clusters, where clusters are added in decreasing order of number of tasks per cluster.

Figure 5 shows these results. As expected, we observe that average performance across the three held-out clusters improves as we add additional clusters and tasks to instruction tuning (with the exception of the sentiment analysis cluster), confirming the benefits of our proposed instruction tuning approach on zero-shot performance on novel tasks. It is further interesting to see that, for the seven clusters we test, the performance does not appear to saturate, implying that performance may further improve with even more clusters added to instruction tuning. Of note, this ablation does not allow us to draw conclusions about which instruction tuning cluster contributes the most to each evaluation cluster, although we see minimal added value from the sentiment analysis cluster.

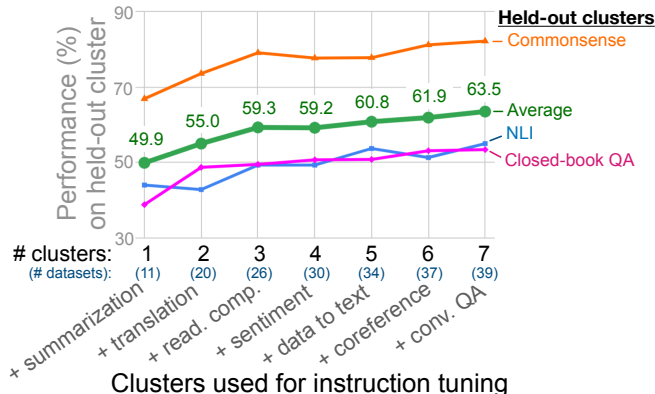


Figure 5: Adding additional task clusters to instruction tuning improves zero-shot performance on held-out task clusters. The evaluation tasks are the following. Commonsense: CoPA, HellaSwag, PiQA, and StoryCloze. NLI: ANLI R1–R3, QNLI, RTE, SNLI, and WNLI. Closed-book QA: ARC easy, ARC challenge, Natural Questions, and TriviaQA.

4.2 SCALING LAWS

As Brown et al. (2020) shows that zero and few-shot capabilities of language models substantially improve for larger models, we next explore how the benefits of instruction tuning are affected by model scale. Using the same cluster split as in the previous ablation study, we evaluate the effect of instruction tuning on models of size 422M, 2B, 8B, 68B, and 137B parameters.

Figure 6 shows these results. First, as a check of correctness, we confirm in Figure 6A that, for all model sizes, performance on tasks that were seen during instruction tuning improves by a large margin over the untuned model without instruction tuning, as would be highly expected. As for held-out tasks, in Figure 6B we see that for the two models on the order of 100B parameters, instruction tuning substantially improves performance, as is also expected given the prior results in our paper. The behavior on held-out tasks for the 8B and smaller models, however, is thought-provoking—instruction tuning actually hurts performance on held-out tasks. One potential explanation for this result could be that for small-scale models, learning the ~ 40 tasks used during instruction tuning fills the entire model capacity, causing these models to perform worse on new tasks. For the larger scale models, instruction tuning fills up some model capacity but also teaches these models the ability to follow instructions, allowing them to generalize to new tasks with the remaining capacity.

²We do not use the paraphrase or reading comprehension with commonsense clusters for instruction tuning in this ablation because they are too similar to NLI and commonsense reasoning, respectively.




Figure 6: (A) Performance on tasks seen during instruction tuning improves for all model sizes. (B) The effect of instruction tuning on performance on unseen tasks depends on the model scale. Whereas instruction tuning helps large models generalize to new tasks, for small models it actually hurts generalization to unseen tasks, potentially because all model capacity is used to learn the mixture of instruction tuning tasks.

4.3 INSTRUCTIONS WITH FEW-SHOT EXEMPLARS

So far, we have focused on instruction tuning in the zero-shot setting. Here, we study how instruction tuning can be used when few-shot exemplars are available at inference time. The format for the few-shot setting builds on the zero-shot format. For some input x and output y , let $\text{instruct}(x)$ denote the zero-shot instructions. Then, given k few-shot exemplars $(x_i, y_i)_{i=1}^k$ and a new input x , the instruction format for the few-shot setting is “ $\text{instruct}(x_1) \oplus y_1 \oplus \text{instruct}(x_2) \oplus y_2 \oplus \dots \oplus \text{instruct}(x_k) \oplus y_k \oplus \text{instruct}(x)$ ”, where \oplus denotes string concatenation with a delimiter token inserted in between. At both training and inference time, exemplars are randomly drawn from the training set, and the number of exemplars is capped at 16 and such that the total sequence length is less than 960 tokens. Our experiment uses the same task splits and evaluation procedure as §3, such that few-shot exemplars for an unseen task are only used at inference time.

As shown in Figure 7, few-shot exemplars improve the performance on all task clusters, compared with zero-shot FLAN. Exemplars are especially effective for tasks with large/complex output spaces, such as struct to text, translation, and closed-book QA, potentially because exemplars help the model better understand the output format. In addition, for all task clusters, standard deviation among templates is lower for few-shot FLAN, indicating reduced sensitivity to prompt engineering.




Figure 7: Adding few-shot exemplars to FLAN is a complementary method for improving the performance of instruction-tuned models. The orange bars indicate standard deviation among templates, averaged at the dataset level for each task cluster.

5 RELATED WORK

Our work relates to several broad research areas including zero-shot learning, prompting, multi-task learning, and language models for NLP applications (Radford et al., 2019; Raffel et al., 2020; Brown et al., 2020; Efrat & Levy, 2020; Aghajanyan et al., 2021; Li & Liang, 2021, *inter alia*). We describe prior work for these broad areas in an extended related work section (Appendix E), and here we describe two subareas narrower in scope that perhaps relate most closely to our work.

The way we ask a model to respond to instructions is similar to QA-based task formulation (Kumar et al., 2016; McCann et al., 2018), which aims to unify NLP tasks by casting them as QA over a context. Though these methods are very similar to ours, they mostly focus on multi-task learning instead of zero-shot learning, and—as noted by Liu et al. (2021)—they are generally not motivated by using existing knowledge in pretrained LMs. Moreover, our work supercedes recent work such as Chai et al. (2020) and Zhong et al. (2021) in terms of both model scale and scope of tasks.

The success of language models has led to nascent research on the ability of models to follow instructions. Most recently, Mishra et al. (2021) finetune 140M parameter BART on instructions with few-shot exemplars, and evaluate its few-shot abilities on unseen tasks—this is very similar to our few-shot instruction tuning result from §4.3. This promising result (as well as one from Ye et al. (2021), which does not emphasize instructions as much) suggests that finetuning on a collection of tasks improves few-shot performance on unseen tasks, even at a smaller model scale. At a model scale similar to ours, OpenAI has a production version of GPT-3 called the “Instruct Series”, which, although they have not yet released a research paper, appears to have goals similar to our work.³

6 DISCUSSION

Our paper has explored a simple question in zero-shot prompting: does finetuning a model on a collection of tasks phrased as instructions improve its performance on unseen tasks? We operationalize this question via instruction tuning, a simple method that combines appealing aspects of both the pretrain–finetune and prompting paradigms. Our instruction-tuned model, FLAN, improves performance against an untuned model and surpasses zero-shot GPT-3 on the majority of tasks that we evaluate on. Ablation studies reveal that performance on unseen tasks improves with the number of instruction tuning task clusters, and, interestingly, that performance improvements from instruction tuning emerge only with sufficient model scale.

The diverse capabilities of language models at scale have drawn attention to the tradeoffs between specialist models (one model per task) and generalist models (one model for many tasks; Arivazhagan et al., 2019; Pratap et al., 2020), for which our study has potential implications. Although one might expect labeled data to have the most natural role in improving specialist models, instruction tuning demonstrates how labeled data can be used to help large language models perform many, unseen tasks. In other words, the positive effect of instruction tuning on cross-task generalization shows that task-specific training is complementary to general language modeling and motivates further research on generalist models.

As for limitations of our study, there is a degree of subjectivity in assigning tasks to clusters (though we try to use accepted categorizations in the literature), and we only explore the use of relatively short instructions of typically a single sentence (c.f. detailed instructions given to crowd-workers). A limitation for our evaluation is that individual examples might have appeared in the models’ pretraining data, which includes web documents, though in post-hoc analysis (Appendix D) we do not find any evidence that data overlap substantially impacted the results. Finally, the scale of FLAN 137B makes it costly to serve. Future work on instruction tuning could include gathering/generating even more task clusters for finetuning, cross-lingual experiments, using FLAN to generate data for training downstream classifiers, and using finetuning to improve model behavior with respect to bias and fairness (Solaiman & Dennison, 2021).

7 CONCLUSIONS

This paper has explored a simple method for improving the ability of language models at scale to perform zero-shot tasks based on purely on instructions. Our instruction-tuned model, FLAN, compares favorably against GPT-3 and signals the potential ability for language models at scale to follow instructions. We hope that our paper will spur further research on instructions-based NLP, zero-shot learning, and using labeled data to improve large language models.

³<https://beta.openai.com/docs/engines/instruct-series-beta>

ETHICAL CONSIDERATIONS

This work uses language models, for which the risks and potential harms are discussed in Bender & Koller (2020), Brown et al. (2020), Bender et al. (2021), Patterson et al., (2021), and others. As our contribution in this paper is not a pretrained language model itself but rather an empirical study of how instruction tuning affects the zero-shot performance of a language model on unseen tasks, we additionally highlight several relevant ethical considerations. First, labeled datasets such as those we use for finetuning can contain undesirable biases, and these biases can be propagated into zero-shot applications of the model on downstream tasks. Next, instruction-tuned models can potentially require less data and expertise to use; such lower barriers to access could increase both the benefits and associated risks of such models. Finally, we want to be clear that, for most datasets, supervised models such as BERT and T5 still outperform the zero-shot instructions-only model in this paper.

ENVIRONMENTAL CONSIDERATIONS

We use the same pretrained language models as Austin et al. (2021). The energy cost and carbon footprint for the pretrained models were 451 MWh and 26 tCO₂e, respectively. The additional instruction tuning gradient-steps for finetuning FLAN is less than 2% of the number of pretraining steps, and so the estimated additional energy cost is comparatively smaller.

AUTHOR CONTRIBUTIONS

Maarten Bosma conceived the original idea and implemented the first version of FLAN. Vincent Zhao prototyped the training and evaluation pipelines, as well as rank classification. Kelvin Guu proposed and implemented the idea of task clusters and evaluation using inter-cluster splits. Jason Wei, Maarten Bosma, Vincent Zhao, and Adams Wei Yu implemented the NLP tasks. Jason Wei, Vincent Zhao, and Adams Wei Yu conducted and managed most of the experiments. Jason Wei designed and ran the ablation studies. Jason Wei, Maarten Bosma, and Quoc V. Le wrote most of the paper. Jason Wei, Maarten Bosma, and Nan Du obtained the zero and few-shot baselines. Vincent Zhao and Kelvin Guu designed, implemented, and conducted the few-shot FLAN experiments. Maarten Bosma and Jason Wei ran the data contamination analysis. Brian Lester ran the prompt tuning experiments. Quoc V. Le and Andrew M. Dai advised, provided high-level guidance and helped edit the paper.

ACKNOWLEDGEMENTS

We thank Ed Chi, Slav Petrov, Yuchen Lin, Dan Garrette, Ruibo Liu, and Clara Meister for providing feedback on our manuscript. We thank Adam Roberts, Liam Fedus, Hyung Won Chung, and Noam Shazeer for helping debug some of our models. We thank Ellie Pavlick for feedback on the study design during the middle stages of the project. We thank Daniel De Freitas Adiwardana for helping initiate the project, large language model advising, and giving us access to some computational resources. Finally, we thank the team involved in pretraining of the Base LM model: Daniel De Freitas Adiwardana, Noam Shazeer, Yanping Huang, Dmitry Lepikhin, Dehao Chen, Yuanzhong Xu and Zhifeng Chen.

REFERENCES

- Armen Aghajanyan, Anchit Gupta, Akshat Shrivastava, Xilun Chen, Luke Zettlemoyer, and Sonal Gupta. Muppet: Massive multi-task representations with pre-finetuning. *arXiv preprint arXiv:2101.11038*, 2021. URL <https://arxiv.org/abs/2101.11038>.
- Naveen Arivazhagan, Ankur Bapna, Orhan Firat, Dmitry Lepikhin, Melvin Johnson, Maxim Krikun, Mia Xu Chen, Yuan Cao, George Foster, Colin Cherry, et al. Massively multilingual neural machine translation in the wild: Findings and challenges. *arXiv preprint arXiv:1907.05019*, 2019. URL <https://arxiv.org/abs/1907.05019>.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, and Charles Sutton. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021. URL <https://arxiv.org/abs/2108.07732>.
- Amittai Axelrod, Xiaodong He, and Jianfeng Gao. Domain adaptation via pseudo in-domain data selection. In *Proceedings of the 2011 Conference on Empirical Methods in Natural Language Processing*, pp. 355–362, 2011. URL <https://aclanthology.org/D11-1033>.
- Marta Bañón, Pinzhen Chen, Barry Haddow, Kenneth Heafield, Hieu Hoang, Miquel Esplà-Gomis, Mikel L. Forcada, Amir Kamran, Faheem Kirefu, Philipp Koehn, Sergio Ortiz Rojas, Leopoldo Pla Sempere, Gema Ramírez-Sánchez, Elsa Sarrias, Marek Strelec, Brian Thompson, William Waites, Dion Wiggins, and Jaume Zaragoza. ParaCrawl: Web-scale acquisition of parallel corpora. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 4555–4567, 2020. URL <https://aclanthology.org/2020.acl-main.417>.
- Emily M. Bender and Alexander Koller. Climbing towards NLU: On meaning, form, and understanding in the age of data. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 5185–5198, 2020. URL <https://aclanthology.org/2020.acl-main.463>.
- Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency, FAccT ’21*, pp. 610–623. Association for Computing Machinery, 2021. URL <https://doi.org/10.1145/3442188.3445922>.
- Luisa Bentivogli, Peter Clark, Ido Dagan, and Danilo Giampiccolo. The Fifth PASCAL Recognizing Textual Entailment Challenge. In *TAC*, 2009. URL <https://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.232.1231&rep=rep1&type=pdf>.
- Yonatan Bisk, Rowan Zellers, Ronan Le Bras, Jianfeng Gao, and Yejin Choi. PIQA: Reasoning about physical commonsense in natural language. In *Thirty-Fourth AAAI Conference on Artificial Intelligence*, 2020. URL <https://arxiv.org/abs/1911.11641>.
- Ondřej Bojar, Christian Buck, Christian Federmann, Barry Haddow, Philipp Koehn, Christof Monz, Matt Post, and Lucia Specia (eds.). *Proceedings of the Ninth Workshop on Statistical Machine Translation*, 2014. URL <https://aclanthology.org/W14-3300>.
- Ondřej Bojar, Christian Buck, Rajen Chatterjee, Christian Federmann, Liane Guillou, Barry Haddow, Matthias Huck, Antonio Jimeno Yepes, Aurélie Névél, Mariana Neves, Pavel Pecina, Martin Popel, Philipp Koehn, Christof Monz, Matteo Negri, Matt Post, Lucia Specia, Karin Verspoor, Jörg Tiedemann, and Marco Turchi (eds.). *Proceedings of the First Conference on Machine Translation: Volume 1, Research Papers*, 2016. URL <https://aclanthology.org/W16-2200>.
- Rishi Bommasani, Drew A. Hudson, E. Adeli, R. Altman, Simran Arora, Sydney von Arx, Michael S. Bernstein, Jeannette Bohg, Antoine Bosselut, Emma Brunskill, E. Brynjolfsson, S. Buch, D. Card, Rodrigo Castellon, Niladri S. Chatterji, Annie Chen, Kathleen Creel, Jared Quincy Davis, Dora Demszky, Chris Donahue, Moussa Doumbouya, Esin Durmus, S. Ermon, J. Etchemendy, Kawin Ethayarajh, L. Fei-Fei, Chelsea Finn, Trevor Gale, Lauren E. Gillespie, Karan Goel, Noah D. Goodman, S. Grossman, Neel Guha, Tatsunori Hashimoto, Peter Henderson, John Hewitt, Daniel E. Ho,

- Jenny Hong, Kyle Hsu, Jing Huang, Thomas F. Icard, Saahil Jain, Dan Jurafsky, Pratyusha Kalluri, Siddharth Karamcheti, G. Keeling, Fereshte Khani, O. Khattab, Pang Wei Koh, M. Krass, Ranjay Krishna, Rohith Kuditipudi, Ananya Kumar, Faisal Ladhak, Mina Lee, Tony Lee, J. Leskovec, Isabelle Levent, Xiang Lisa Li, Xuechen Li, Tengyu Ma, Ali Malik, Christopher D. Manning, Suvir P. Mirchandani, Eric Mitchell, Zanele Munyikwa, Suraj Nair, Avani Narayan, D. Narayanan, Ben Newman, Allen Nie, J. C. Niebles, H. Nilforoshan, Julian Nyarko, Giray Ogut, Laurel Orr, Isabel Papadimitriou, Joon Sung Park, C. Piech, Eva Portelance, Christopher Potts, Aditi Raghunathan, Robert Reich, Hongyu Ren, Frieda Rong, Yusuf H. Roohani, Camilo Ruiz, Jack Ryan, Christopher R’e, D. Sadigh, Shiori Sagawa, Keshav Santhanam, Andy Shih, K. Srinivasan, Alex Tamkin, Rohan Taori, Armin W. Thomas, Florian Tramèr, Rose E. Wang, William Wang, Bohan Wu, Jiajun Wu, Yuhuai Wu, Sang Michael Xie, Michihiro Yasunaga, Jiaxuan You, M. Zaharia, Michael Zhang, Tianyi Zhang, Xikun Zhang, Yuhui Zhang, Lucia Zheng, Kaitlyn Zhou, and Percy Liang. On the opportunities and risks of foundation models. 2021. URL <https://arxiv.org/abs/2108.07258>.
- Samuel R. Bowman, Gabor Angeli, Christopher Potts, and Christopher D. Manning. A large annotated corpus for learning natural language inference. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pp. 632–642, 2015. URL <https://aclanthology.org/D15-1075>.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, volume 33, pp. 1877–1901, 2020. URL <https://proceedings.neurips.cc/paper/2020/file/1457c0d6bfbcb4967418bfb8ac142f64a-Paper.pdf>.
- Duo Chai, Wei Wu, Qinghong Han, Fei Wu, and Jiwei Li. Description based text classification with reinforcement learning. In *Proceedings of the International Conference on Machine Learning*, pp. 1371–1382. PMLR, 2020. URL <http://proceedings.mlr.press/v119/chai20a/chai20a.pdf>.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde, Jared Kaplan, Harri Edwards, Yura Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021. URL <https://arxiv.org/abs/2107.03374>.
- Eunsol Choi, He He, Mohit Iyyer, Mark Yatskar, Wen-tau Yih, Yejin Choi, Percy Liang, and Luke Zettlemoyer. QuAC: Question answering in context. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 2174–2184, 2018. URL <https://aclanthology.org/D18-1241>.
- Christopher Clark, Kenton Lee, Ming-Wei Chang, Tom Kwiatkowski, Michael Collins, and Kristina Toutanova. BoolQ: Exploring the surprising difficulty of natural yes/no questions. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pp. 2924–2936, 2019a. URL <https://aclanthology.org/N19-1300>.
- Kevin Clark, Minh-Thang Luong, Urvashi Khandelwal, Christopher D. Manning, and Quoc V. Le. BAM! born-again multi-task networks for natural language understanding. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 5931–5937, 2019b. URL <https://aclanthology.org/P19-1595>.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? Try ARC, the AI2 reasoning challenge. *arXiv preprint arXiv:1803.05457*, 2018. URL <https://arxiv.org/abs/1803.05457>.
- Ronan Collobert, Jason Weston, Léon Bottou, Michael Karlen, Koray Kavukcuoglu, and Pavel Kuksa. Natural language processing (almost) from scratch. *Journal of Machine Learning*

- Research*, 12:2493–2537, 2011. URL <https://www.jmlr.org/papers/volume12/collobert11a/collobert11a.pdf>.
- Michele Corazza, Stefano Menini, Elena Cabrio, Sara Tonelli, and Serena Villata. Hybrid emoji-based masked language models for zero-shot abusive language detection. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pp. 943–949, 2020. URL <https://aclanthology.org/2020.findings-emnlp.84>.
- Ido Dagan, Oren Glickman, and Bernardo Magnini. The PASCAL Recognising Textual Entailment challenge. In *Proceedings of the First International Conference on Machine Learning Challenges: Evaluating Predictive Uncertainty Visual Object Classification, and Recognizing Textual Entailment*, MLCW’05, pp. 177–190, 2005. URL https://doi.org/10.1007/11736790_9.
- Andrew M Dai and Quoc V Le. Semi-supervised sequence learning. In *Proceedings of the Conference on Neural Information Processing Systems*, 2015. URL <https://papers.nips.cc/paper/2015/file/7137debd45ae4d0ab9aa953017286b20-Paper.pdf>.
- Marie-Catherine De Marneffe, Mandy Simons, and Judith Tonhauser. The CommitmentBank: Investigating projection in naturally occurring discourse. In *Proceedings of Sinn und Bedeutung*, volume 23, pp. 107–124, 2019. URL <https://ojs.ub.uni-konstanz.de/sub/index.php/sub/article/view/601>.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pp. 4171–4186, 2019. URL <https://aclanthology.org/N19-1423>.
- William B. Dolan and Chris Brockett. Automatically constructing a corpus of sentential paraphrases. In *Proceedings of the Third International Workshop on Paraphrasing (IWP2005)*, 2005. URL <https://aclanthology.org/I05-5002>.
- Dheeru Dua, Yizhong Wang, Pradeep Dasigi, Gabriel Stanovsky, Sameer Singh, and Matt Gardner. DROP: A reading comprehension benchmark requiring discrete reasoning over paragraphs. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pp. 2368–2378, 2019. URL <https://aclanthology.org/N19-1246>.
- Nadir Durrani, Barry Haddow, Philipp Koehn, and Kenneth Heafield. Edinburgh’s phrase-based machine translation systems for WMT-14. In *Proceedings of the Ninth Workshop on Statistical Machine Translation*, pp. 97–104, 2014. URL <https://aclanthology.org/W14-3309>.
- Ondřej Dušek, David M. Howcroft, and Verena Rieser. Semantic noise matters for neural natural language generation. In *Proceedings of the 12th International Conference on Natural Language Generation*, pp. 421–426, 2019. URL <https://aclanthology.org/W19-8652>.
- Sergey Edunov, Myle Ott, Michael Auli, and David Grangier. Understanding back-translation at scale. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 489–500, 2018. URL <https://aclanthology.org/D18-1045>.
- Avia Efrat and Omer Levy. The Turking Test: Can language models understand instructions? *arXiv preprint arXiv:2010.11982*, 2020. URL <https://arxiv.org/abs/2010.11982>.
- Alexander Fabbri, Irene Li, Tianwei She, Suyi Li, and Dragomir Radev. Multi-news: A large-scale multi-document summarization dataset and abstractive hierarchical model. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 1074–1084, 2019. URL <https://aclanthology.org/P19-1102>.
- Fast.AI. Yelp Sentiment Classification Dataset. <https://course.fast.ai/datasets>.
- William Fedus, Barret Zoph, and Noam Shazeer. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *arXiv preprint arXiv:2101.03961*, 2021. URL <https://arxiv.org/abs/2101.03961>.

- Chelsea Finn, Pieter Abbeel, and Sergey Levine. Model-agnostic meta-learning for fast adaptation of deep networks. In *Proceedings of the International Conference on Machine Learning (ICML)*, pp. 1126–1135, 2017. URL <https://arxiv.org/abs/1703.03400>.
- Tianyu Gao, Adam Fisch, and Danqi Chen. Making pre-trained language models better few-shot learners. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pp. 3816–3830, 2021. URL <https://aclanthology.org/2021.acl-long.295>.
- Claire Gardent, Anastasia Shimorina, Shashi Narayan, and Laura Perez-Beltrachini. The WebNLG challenge: Generating text from RDF data. In *Proceedings of the 10th International Conference on Natural Language Generation*, pp. 124–133, 2017. URL <https://aclanthology.org/W17-3518>.
- Sebastian Gehrmann, Tosin Adewumi, Karmanya Aggarwal, Pawan Sasanka Ammanamanchi, Anuoluwapo Aremu, Antoine Bosselut, Khyathi Raghavi Chandu, Miruna-Adriana Clinciu, Dipanjan Das, Kaustubh Dhole, Wanyu Du, Esin Durmus, Ondřej Dušek, Chris Chinenye Emezue, Varun Gangal, Cristina Garbacea, Tatsunori Hashimoto, Yufang Hou, Yacine Jernite, Harsh Jhamtani, Yangfeng Ji, Shailza Jolly, Mihir Kale, Dhruv Kumar, Faisal Ladhak, Aman Madaan, Mounica Maddela, Khyati Mahajan, Saad Mahamood, Bodhisattwa Prasad Majumder, Pedro Henrique Martins, Angelina McMillan-Major, Simon Mille, Emiel van Miltenburg, Moin Nadeem, Shashi Narayan, Vitaly Nikolaev, Andre Niyongabo Rubungo, Salomey Osei, Ankur Parikh, Laura Perez-Beltrachini, Niranjan Ramesh Rao, Vikas Raunak, Juan Diego Rodriguez, Sashank Santhanam, João Sedoc, Thibault Sellam, Samira Shaikh, Anastasia Shimorina, Marco Antonio Sobrevilla Cabezudo, Hendrik Strobelt, Nishant Subramani, Wei Xu, Diyi Yang, Akhila Yerukola, and Jiawei Zhou. The GEM benchmark: Natural language generation, its evaluation and metrics. In *Proceedings of the 1st Workshop on Natural Language Generation, Evaluation, and Metrics (GEM 2021)*, pp. 96–120, 2021. URL <https://aclanthology.org/2021.gem-1.10>.
- Danilo Giampiccolo, Bernardo Magnini, Ido Dagan, and Bill Dolan. The third PASCAL recognizing textual entailment challenge. In *Proceedings of the ACL-PASCAL Workshop on Textual Entailment and Paraphrasing*, pp. 1–9, 2007. URL <https://aclanthology.org/W07-1401>.
- Bogdan Gliwa, Iwona Mochol, Maciej Biesek, and Aleksander Wawer. SAMSum corpus: A human-annotated dialogue dataset for abstractive summarization. In *Proceedings of the 2nd Workshop on New Frontiers in Summarization*, pp. 70–79, 2019. URL <https://aclanthology.org/D19-5409>.
- Alec Go, Richa Bhayani, and Lei Huang. Twitter sentiment classification using distant supervision. *CS224N project report, Stanford*, 1(12):2009, 2009. URL <https://www-cs.stanford.edu/people/alecmgo/papers/TwitterDistantSupervision09.pdf>.
- Dan Goldwasser and Dan Roth. Learning from natural instructions. *Machine learning*, 94(2):205–232, 2014. URL <https://link.springer.com/article/10.1007/s10994-013-5407-y>.
- Max Grusky, Mor Naaman, and Yoav Artzi. Newsroom: A dataset of 1.3 million summaries with diverse extractive strategies. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pp. 708–719, 2018. URL <https://aclanthology.org/N18-1065>.
- R Bar Haim, Ido Dagan, Bill Dolan, Lisa Ferro, Danilo Giampiccolo, Bernardo Magnini, and Idan Szpektor. The Second PASCAL Recognising Textual Entailment Challenge. In *Proceedings of the Second PASCAL Challenges Workshop on Recognising Textual Entailment*, 2006. URL <http://www.cs.biu.ac.il/~szpekti/papers/RTE2-organizers.pdf>.
- Luheng He, Mike Lewis, and Luke Zettlemoyer. Question-answer driven semantic role labeling: Using natural language to annotate natural language. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pp. 643–653, 2015. URL <https://aclanthology.org/D15-1076>.

- Eduard Hovy, Laurie Gerber, Ulf Hermjakob, Chin-Yew Lin, and Deepak Ravichandran. Toward semantics-based answer pinpointing. In *Proceedings of the First International Conference on Human Language Technology Research*, 2001. URL <https://www.aclweb.org/anthology/H01-1069>.
- Jeremy Howard and Sebastian Ruder. Universal language model fine-tuning for text classification. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 328–339, 2018. URL <https://aclanthology.org/P18-1031>.
- Lifu Huang, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. Cosmos QA: Machine reading comprehension with contextual commonsense reasoning. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pp. 2391–2401, 2019. URL <https://aclanthology.org/D19-1243>.
- Melvin Johnson, Mike Schuster, Quoc V. Le, Maxim Krikun, Yonghui Wu, Zhifeng Chen, Nikhil Thorat, Fernanda Viégas, Martin Wattenberg, Greg Corrado, Macduff Hughes, and Jeffrey Dean. Google’s multilingual neural machine translation system: Enabling zero-shot translation. *Transactions of the Association for Computational Linguistics*, 5:339–351, 2017. URL <https://aclanthology.org/Q17-1024>.
- Mandar Joshi, Eunsol Choi, Daniel Weld, and Luke Zettlemoyer. TriviaQA: A large scale distantly supervised challenge dataset for reading comprehension. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 1601–1611, 2017. URL <https://aclanthology.org/P17-1147>.
- Daniel Khashabi, Snigdha Chaturvedi, Michael Roth, Shyam Upadhyay, and Dan Roth. Looking beyond the surface: A challenge set for reading comprehension over multiple sentences. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pp. 252–262, 2018. URL <https://aclanthology.org/N18-1023>.
- Dimitrios Kotzias, Misha Denil, Nando de Freitas, and Padhraic Smyth. From group to individual labels using deep features. *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2015. URL <https://dl.acm.org/doi/10.1145/2783258.2783380>.
- Taku Kudo and John Richardson. Sentencepiece: A simple and language independent subword tokenizer and detokenizer for neural text processing. In Eduardo Blanco and Wei Lu (eds.), *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing, EMNLP 2018: System Demonstrations, Brussels, Belgium, October 31 - November 4, 2018*, pp. 66–71. Association for Computational Linguistics, 2018. doi: 10.18653/v1/d18-2012. URL <https://doi.org/10.18653/v1/d18-2012>.
- Ankit Kumar, Ozan Irsoy, Peter Ondruska, Mohit Iyyer, James Bradbury, Ishaan Gulrajani, Victor Zhong, Romain Paulus, and Richard Socher. Ask me anything: Dynamic memory networks for natural language processing. In *Proceedings of the International Conference on Machine Learning*, pp. 1378–1387. PMLR, 2016. URL <https://arxiv.org/abs/1506.07285>.
- Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, Kristina Toutanova, Llion Jones, Matthew Kelcey, Ming-Wei Chang, Andrew M. Dai, Jakob Uszkoreit, Quoc Le, and Slav Petrov. Natural Questions: A benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:452–466, 2019. URL <https://aclanthology.org/Q19-1026>.
- Faisal Ladhak, Esin Durmus, Claire Cardie, and Kathleen McKeown. WikiLingua: A new benchmark dataset for cross-lingual abstractive summarization. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pp. 4034–4048, 2020. URL <https://aclanthology.org/2020.findings-emnlp.360>.

- Christoph H Lampert, Hannes Nickisch, and Stefan Harmeling. Learning to detect unseen object classes by between-class attribute transfer. In *2009 IEEE Conference on Computer Vision and Pattern Recognition*, pp. 951–958. IEEE, 2009. URL <https://ieeexplore.ieee.org/document/5206594>.
- Anne Lauscher, Vinit Ravishankar, Ivan Vulić, and Goran Glavaš. From zero to hero: On the limitations of zero-shot language transfer with multilingual Transformers. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 4483–4499, 2020. URL <https://aclanthology.org/2020.emnlp-main.363>.
- Kenton Lee, Ming-Wei Chang, and Kristina Toutanova. Latent retrieval for weakly supervised open domain question answering. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 6086–6096, 2019. URL <https://aclanthology.org/P19-1612>.
- Dmitry Lepikhin, HyounJoong Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping Huang, Maxim Krikun, Noam Shazeer, and Zhifeng Chen. Gshard: Scaling giant models with conditional computation and automatic sharding. In *International Conference on Learning Representations*, 2020. URL <https://openreview.net/forum?id=qrwe7XHTmYb>.
- Brian Lester, Rami Al-Rfou, and Noah Constant. The power of scale for parameter-efficient prompt tuning. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, 2021. URL <https://arxiv.org/abs/2104.08691>.
- Hector Levesque, Ernest Davis, and Leora Morgenstern. The Winograd Schema Challenge. In *Thirteenth International Conference on the Principles of Knowledge Representation and Reasoning*, 2012. URL <https://dl.acm.org/doi/10.5555/3031843.3031909>.
- Omer Levy, Minjoon Seo, Eunsol Choi, and Luke Zettlemoyer. Zero-shot relation extraction via reading comprehension. In *Proceedings of the 21st Conference on Computational Natural Language Learning (CoNLL 2017)*, pp. 333–342, 2017. URL <https://aclanthology.org/K17-1034>.
- Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Veselin Stoyanov, and Luke Zettlemoyer. BART: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 7871–7880, 2020. URL <https://aclanthology.org/2020.acl-main.703>.
- Xiang Lisa Li and Percy Liang. Prefix-tuning: Optimizing continuous prompts for generation. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pp. 4582–4597, 2021. URL <https://aclanthology.org/2021.acl-long.353>.
- Xiaoya Li, Jingrong Feng, Yuxian Meng, Qinghong Han, Fei Wu, and Jiwei Li. A unified MRC framework for named entity recognition. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 5849–5859, 2020. URL <https://aclanthology.org/2020.acl-main.519>.
- Xin Li and Dan Roth. Learning question classifiers. In *COLING 2002: The 19th International Conference on Computational Linguistics*, 2002. URL <https://www.aclweb.org/anthology/C02-1150>.
- Bill Yuchen Lin, Wangchunshu Zhou, Ming Shen, Pei Zhou, Chandra Bhagavatula, Yejin Choi, and Xiang Ren. CommonGen: A constrained text generation challenge for generative commonsense reasoning. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pp. 1823–1840, 2020. URL <https://aclanthology.org/2020.findings-emnlp.165>.
- Han Liu, Xiaotong Zhang, Lu Fan, Xuandi Fu, Qimai Li, Xiao-Ming Wu, and Albert Y.S. Lam. Reconstructing capsule networks for zero-shot intent classification. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pp. 4799–4809, 2019a. URL <https://aclanthology.org/D19-1486>.

- Pengfei Liu, Weizhe Yuan, Jinlan Fu, Zhenbao Jiang, Hiroaki Hayashi, and Graham Neubig. Pre-train, prompt, and predict: A systematic survey of prompting methods in natural language processing. *arXiv preprint arXiv:2107.13586*, 2021. URL <https://arxiv.org/abs/2107.13586>.
- Xiaodong Liu, Pengcheng He, Weizhu Chen, and Jianfeng Gao. Multi-task deep neural networks for natural language understanding. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 4487–4496, 2019b. URL <https://aclanthology.org/P19-1441>.
- Yinhan Liu, Jiatao Gu, Naman Goyal, Xian Li, Sergey Edunov, Marjan Ghazvininejad, Mike Lewis, and Luke Zettlemoyer. Multilingual denoising pre-training for neural machine translation. *Transactions of the Association for Computational Linguistics*, 8:726–742, 2020. URL <https://aclanthology.org/2020.tacl-1.47>.
- Minh-Thang Luong, Quoc V Le, Ilya Sutskever, Oriol Vinyals, and Lukasz Kaiser. Multi-task sequence to sequence learning. *Proceedings of ICLR*, 2016. URL <https://arxiv.org/abs/1511.06114>.
- Andrew L. Maas, Raymond E. Daly, Peter T. Pham, Dan Huang, Andrew Y. Ng, and Christopher Potts. Learning word vectors for sentiment analysis. In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*, pp. 142–150, Portland, Oregon, USA, June 2011. Association for Computational Linguistics. URL <http://www.aclweb.org/anthology/P11-1015>.
- Bryan McCann, Nitish Shirish Keskar, Caiming Xiong, and Richard Socher. The natural language decathlon: Multitask learning as question answering. 2018. URL <https://arxiv.org/abs/1806.08730>.
- John McCarthy. *Programs with common sense*. RLE and MIT computation center, 1960. URL <http://jmc.stanford.edu/articles/mcc59/mcc59.pdf>.
- Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. Can a suit of armor conduct electricity? A new dataset for open book question answering. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 2381–2391, 2018. URL <https://aclanthology.org/D18-1260>.
- Swaroop Mishra, Daniel Khashabi, Chitta Baral, and Hannaneh Hajishirzi. Natural Instructions: Benchmarking generalization to new tasks from natural language instructions. *arXiv preprint arXiv:2104.08773*, 2021. URL <https://arxiv.org/abs/2104.08773>.
- Nasrin Mostafazadeh, Nathanael Chambers, Xiaodong He, Devi Parikh, Dhruv Batra, Lucy Vanderwende, Pushmeet Kohli, and James Allen. A corpus and cloze evaluation for deeper understanding of commonsense stories. In *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pp. 839–849, 2016. URL <https://aclanthology.org/N16-1098>.
- Linyong Nan, Dragomir Radev, Rui Zhang, Amrit Rau, Abhinand Sivaprasad, Chiachun Hsieh, Xiangru Tang, Aadit Vyas, Neha Verma, Pranav Krishna, Yangxiaokang Liu, Nadia Irwanto, Jessica Pan, Faiaz Rahman, Ahmad Zaidi, Mutethia Mutuma, Yasin Tarabar, Ankit Gupta, Tao Yu, Yi Chern Tan, Xi Victoria Lin, Caiming Xiong, Richard Socher, and Nazneen Fatema Rajani. DART: Open-domain structured data record to text generation. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pp. 432–447, 2021. URL <https://aclanthology.org/2021.naacl-main.37>.
- Courtney Napoles, Matthew Gormley, and Benjamin Van Durme. Annotated Gigaword. In *Proceedings of the Joint Workshop on Automatic Knowledge Base Construction and Web-scale Knowledge Extraction (AKBC-WEKEX)*, pp. 95–100, 2012. URL <https://aclanthology.org/W12-3018>.
- Shashi Narayan, Shay B. Cohen, and Mirella Lapata. Don’t give me the details, just the summary! topic-aware convolutional neural networks for extreme summarization. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 1797–1807, 2018. URL <https://aclanthology.org/D18-1206>.

- Yixin Nie, Adina Williams, Emily Dinan, Mohit Bansal, Jason Weston, and Douwe Kiela. Adversarial NLI: A new benchmark for natural language understanding. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 4885–4901, 2020. URL <https://aclanthology.org/2020.acl-main.441>.
- Matthew E. Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. Deep contextualized word representations. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pp. 2227–2237, 2018. URL <https://aclanthology.org/N18-1202>.
- Ngoc-Quan Pham, Jan Niehues, Thanh-Le Ha, and Alexander Waibel. Improving zero-shot translation with language-independent constraints. In *Proceedings of the Fourth Conference on Machine Translation (Volume 1: Research Papers)*, pp. 13–23, 2019. URL <https://aclanthology.org/W19-5202>.
- Mohammad Taher Pilehvar and Jose Camacho-Collados. WiC: the word-in-context dataset for evaluating context-sensitive meaning representations. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pp. 1267–1273, 2019. URL <https://aclanthology.org/N19-1128>.
- Vineel Pratap, Anuroop Sriram, Paden Tomasello, Awni Hannun, Vitaliy Liptchinsky, Gabriel Synnaeve, and Ronan Collobert. Massively multilingual ASR: 50 languages, 1 model, 1 billion parameters. *arXiv preprint arXiv:2007.03001*, 2020. URL <https://arxiv.org/abs/2007.03001>.
- Guanghui Qin and Jason Eisner. Learning how to ask: Querying LMs with mixtures of soft prompts. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*, pp. 5203–5212, 2021. URL <http://cs.jhu.edu/~jason/papers/#qin-eisner-2021>.
- Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. Improving language understanding by generative pre-training. <https://blog.openai.com/language-unsupervised>, 2018.
- Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9, 2019. URL https://d4mucfpsywv.cloudfront.net/better-language-models/language_models_are_unsupervised_multitask_learners.pdf.
- Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020. URL <http://jmlr.org/papers/v21/20-074.html>.
- Altat Rahman and Vincent Ng. Resolving complex cases of definite pronouns: The Winograd schema challenge. In *Proceedings of the 2012 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning*, pp. 777–789, 2012. URL <https://aclanthology.org/D12-1071>.
- Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. SQuAD: 100,000+ questions for machine comprehension of text. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pp. 2383–2392, 2016. URL <https://aclanthology.org/D16-1264>.
- Pranav Rajpurkar, Robin Jia, and Percy Liang. Know what you don’t know: Unanswerable questions for SQuAD. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pp. 784–789, 2018. URL <https://aclanthology.org/P18-2124>.

- Siva Reddy, Danqi Chen, and Christopher D. Manning. CoQA: A conversational question answering challenge. *Transactions of the Association for Computational Linguistics*, 7:249–266, 2019. URL <https://aclanthology.org/Q19-1016>.
- Emily Reif, Daphne Ippolito, Ann Yuan, Andy Coenen, Chris Callison-Burch, and Jason Wei. A recipe for arbitrary text style transfer with large language models. *arXiv preprint arXiv:2109.03910*, 2021. URL <https://arxiv.org/abs/2109.03910>.
- Melissa Roemmele, Cosmin Bejan, and Andrew Gordon. Choice of plausible alternatives: An evaluation of commonsense causal reasoning, 2011. URL <https://www.aaai.org/ocs/index.php/SSS/SSS11/paper/view/2418>.
- Bernardino Romera-Paredes and Philip Torr. An embarrassingly simple approach to zero-shot learning. In *Proceedings of the International Conference on Machine Learning*, pp. 2152–2161, 2015. URL <https://proceedings.mlr.press/v37/romera-paredes15.pdf>.
- Sebastian Ruder. An overview of multi-task learning in deep neural networks. *arXiv preprint arXiv:1706.05098*, 2017. URL <https://arxiv.org/abs/1706.05098>.
- Keisuke Sakaguchi, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. WinoGrande: An adversarial winograd schema challenge at scale. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pp. 8732–8740, 2020. URL <https://arxiv.org/abs/1907.10641>.
- David Saxton, Edward Grefenstette, Felix Hill, and Pushmeet Kohli. Analysing mathematical reasoning abilities of neural models. *Proceedings of the International Conference on Learning Representations*, 2019. URL <https://arxiv.org/pdf/1904.01557>.
- Timo Schick and Hinrich Schütze. Exploiting cloze-questions for few-shot text classification and natural language inference. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pp. 255–269, 2021. URL <https://aclanthology.org/2021.eacl-main.20>.
- Abigail See, Peter J. Liu, and Christopher D. Manning. Get to the point: Summarization with pointer-generator networks. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 1073–1083, 2017. URL <https://aclanthology.org/P17-1099>.
- Rico Sennrich, Barry Haddow, and Alexandra Birch. Edinburgh neural machine translation systems for WMT 16. In *Proceedings of the First Conference on Machine Translation: Volume 2, Shared Task Papers*, pp. 371–376, 2016. URL <https://aclanthology.org/W16-2323>.
- Noam Shazeer and Mitchell Stern. Adafactor: Adaptive learning rates with sublinear memory cost. In *International Conference on Machine Learning*, pp. 4596–4604. PMLR, 2018. URL <https://arxiv.org/abs/1804.04235>.
- Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D. Manning, Andrew Ng, and Christopher Potts. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, pp. 1631–1642, 2013. URL <https://aclanthology.org/D13-1170>.
- Irene Solaiman and Christy Dennison. Process for adapting language models to society (palms) with values-targeted datasets. *arXiv preprint arXiv:2106.10328*, 2021. URL <https://arxiv.org/abs/2106.10328>.
- Shashank Srivastava, Igor Labutov, and Tom Mitchell. Zero-shot learning of classifiers from natural language quantification. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 306–316, 2018. URL <https://aclanthology.org/P18-1029>.
- Derek Tam, Menton Rakesh R., Mohit Bansal, Shashank Srivastava, and Colin Raffel. Improving and simplifying pattern exploiting training. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2021. URL <https://arxiv.org/pdf/2103.11955>.

- Joaquin Vanschoren. Meta-learning: A survey. *arXiv preprint arXiv:1810.03548*, 2018. URL <https://arxiv.org/abs/1810.03548>.
- Marc Velay and Fabrice Daniel. Seq2seq and multi-task learning for joint intent and content extraction for domain specific interpreters. *arXiv preprint arXiv:1808.00423*, 2018. URL <https://arxiv.org/abs/1808.00423>.
- Alex Wang, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel Bowman. GLUE: A multi-task benchmark and analysis platform for natural language understanding. In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pp. 353–355, 2018. URL <https://aclanthology.org/W18-5446>.
- Alex Wang, Yada Pruksachatkun, Nikita Nangia, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel R Bowman. Superglue: A stickier benchmark for general-purpose language understanding systems. *Conference on Neural Information Processing Systems (NeurIPS)*, 2019a. URL <https://arxiv.org/abs/1905.00537>.
- Lu Wang and Wang Ling. Neural network-based abstract generation for opinions and arguments. In *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pp. 47–57, 2016. URL <https://www.aclweb.org/anthology/N16-1007>.
- Yiren Wang, Yingce Xia, Tianyu He, Fei Tian, Tao Qin, ChengXiang Zhai, and Tie-Yan Liu. Multi-agent dual learning. In *Proceedings of the International Conference on Learning Representations (ICLR) 2019*, 2019b. URL <https://openreview.net/forum?id=HyGhN2A5tm>.
- Zirui Wang, Adams Wei Yu, Orhan Firat, and Yuan Cao. Towards zero-label language learning. *arXiv preprint arXiv:2109.09193*, 2021. URL <https://arxiv.org/abs/2109.09193>.
- Alex Warstadt, Amanpreet Singh, and Samuel R. Bowman. Neural network acceptability judgments. *Transactions of the Association for Computational Linguistics*, 7:625–641, 2019. doi: 10.1162/tacl_a_00290. URL <https://aclanthology.org/Q19-1040>.
- Adina Williams, Nikita Nangia, and Samuel Bowman. A broad-coverage challenge corpus for sentence understanding through inference. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pp. 1112–1122, 2018. URL <http://aclweb.org/anthology/N18-1101>.
- Joseph Worsham and J. Kalita. Multi-task learning for natural language processing in the 2020s: where are we going? *arXiv preprint arXiv:2007.16008*, 2020. URL <https://arxiv.org/abs/2007.16008>.
- Jeff Wu, Long Ouyang, Daniel M Ziegler, Nissan Stiennon, Ryan Lowe, Jan Leike, and Paul Christiano. Recursively summarizing books with human feedback. *arXiv preprint arXiv:2109.10862*, 2021. URL <https://arxiv.org/abs/2109.10862>.
- Wei Wu, Fei Wang, Arianna Yuan, Fei Wu, and Jiwei Li. CorefQA: Coreference resolution as query-based span prediction. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 6953–6963, 2020. URL <https://aclanthology.org/2020.acl-main.622>.
- Qinyuan Ye, Bill Yuchen Lin, and Xiang Ren. Crossfit: A few-shot learning challenge for cross-task generalization in NLP. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2021. URL <https://arxiv.org/abs/2104.08835>.
- Wenpeng Yin, Jamaal Hay, and Dan Roth. Benchmarking zero-shot text classification: Datasets, evaluation and entailment approach. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pp. 3914–3923, 2019. URL <https://aclanthology.org/D19-1404>.

- Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. HellaSwag: Can a machine really finish your sentence? In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 4791–4800, 2019. URL <https://aclanthology.org/P19-1472>.
- Rui Zhang and Joel Tetreault. This email could save your life: Introducing the task of email subject line generation. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, 2019. URL <https://aclanthology.org/P19-1043>.
- Sheng Zhang, Xiaodong Liu, Jingjing Liu, Jianfeng Gao, Kevin Duh, and Benjamin Van Durme. Record: Bridging the gap between human and machine commonsense reading comprehension. *CoRR*, abs/1810.12885, 2018. URL <http://arxiv.org/abs/1810.12885>.
- Xiang Zhang, Junbo Zhao, and Yann LeCun. Character-level convolutional networks for text classification. In *Advances in Neural Information Processing Systems*, volume 28, 2015. URL <https://proceedings.neurips.cc/paper/2015/file/250cf8b51c773f3f8dc8b4be867a9a02-Paper.pdf>.
- Yuan Zhang, Jason Baldridge, and Luheng He. PAWS: Paraphrase adversaries from word scrambling. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pp. 1298–1308, 2019. URL <https://aclanthology.org/N19-1131>.
- Ruiqi Zhong, Kristy Lee, Zheng Zhang, and Dan Klein. Meta-tuning language models to answer prompts better. *arXiv preprint arXiv:2104.04670*, 2021. URL <https://arxiv.org/abs/2104.04670>.

A ADDITIONAL RESULTS

A.1 COMMONSENSE REASONING & COREFERENCE RESOLUTION

Results for five commonsense reasoning datasets are shown in Table 4. FLAN outperforms GPT-3 on StoryCloze (Mostafazadeh et al., 2016) and has comparable performance on CoPA (Roemmele et al., 2011) and PiQA (Bisk et al., 2020). For HellaSwag (Zellers et al., 2019) and ReCoRD (Zhang et al., 2018), however, both Base LM and FLAN performs well below GPT-3. Table 4 also shows results for two coreference resolution datasets. FLAN with the best dev template outperforms zero-shot GPT-3 on Winogrande (Sakaguchi et al., 2020). For WSC273 (Levesque et al., 2012), both Base LM and FLAN are well below GPT-3 (we cannot compute the best dev template result for FLAN since WSC273 does not have a validation set).

One observation is that HellaSwag, PiQA, ReCoRD, WSC273, and Winogrande are directly formulated in the original pretraining objective of language modeling. Hence, the result that GPT-3 outperforms FLAN by a large margin is consistent with the intuition that gains from FLAN will be marginal when instructions are not crucial for describing the given task. Moreover, we note a further limitation with FLAN for these five language modeling tasks, including options actually hurts performance, and so the reported results are for rank classification without options.

	COMMONSENSE REASONING					COREFERENCE	
	CoPA acc.	HellaSwag acc.	PiQA acc.	StoryCloze acc.	ReCoRD acc.	WSC273 acc.	Winogrande acc.
Supervised model	94.8 ^a	47.3 ^b	66.8 ^b	89.2 ^b	93.4 ^a	72.2 ^b	93.8 ^a
Base LM 137B zero-shot	90.0	57.0	80.3*	79.5	87.8*	81.0	68.3
· few-shot	89.0	58.8	80.2*	83.7	87.6*	61.5	68.4
GPT-3 175B zero-shot	91.0	78.9	81.0	83.2	90.2	88.3	70.2
· few-shot	92.0	79.3	82.3	87.7	89.0	88.6	77.7
FLAN 137B zero-shot							
- average template	90.6 std=2.0	56.4 std=0.5	80.9* std=0.8	92.2▲4.5 std=1.3	67.8* std=3.0	80.8 std=3.7	67.3 std=2.5
- best dev template	91.0	56.7	80.5*	93.4▲5.7	72.5*	-	71.2↑1.0

Table 4: Results (accuracy in %) for commonsense reasoning and coreference resolution. For FLAN, we report both the average of up to ten templates, as well as the best dev template. ^aT5-11B, ^bBERT-large. The triangle ▲ indicates improvement over few-shot GPT-3. The up-arrow ↑ indicates improvement only over zero-shot GPT-3. *see data contamination (Appendix D).

A.2 ADDITIONAL DATASETS

The main paper and Appendix A.1 focused on datasets for which GPT-3 results are available for comparison; here we report additional results on datasets for which comparisons with GPT-3 were not available. Results for eight NLU task clusters are shown in Table 5 on the next page, and the results for translation and struct to text are shown in Table 6 on the page after that.

We show FLAN’s performance using the best of up to ten instruction templates as well as the template with the best performance on the dev set. For Base LM, we use the templates from Brown et al. (2020), which were optimized for GPT-3, without performing any prompt engineering to optimize them on our model. For simplicity, we use greedy search for all generative tasks (compared with beam search used in Brown et al. (2020)). Unlike GPT-3, which chooses the number of few-shot exemplars k via best dev set performance, for few-shot Base LM we choose the highest k that fits in the context length of 1024 tokens, from $k \in \{1, 3, 5, 10\}$.

For DROP (Dua et al., 2019) and SQuADv2 (Rajpurkar et al., 2018), based on email correspondence with Brown et al. (2020), their definition of zero-shot differs from ours in that they actually use exemplars, but only from the same passage as the inference question (each passage has more than one question). Hence, GPT-3 zero-shot results are not directly comparable with ours for DROP and SQuADv2. We mark these results using the † symbol. Moreover, it is unclear how to parse the end of an answer for these two datasets, and so we use curly bracket delimiters { and }, where we expect } to indicate the end of the answer.

FLAN 137B												
	Random Guess	Supervised Model	Base LM 137B		GPT-3 175B		zero-shot		few-shot		[k]	#t
			zero- shot	few- shot [k]	zero- shot	few- shot [k]	average template	best dev template	average template	best dev template		
NLI												
ANLI R1	33.3	57.4 ^b	39.6	39.0 [5]	34.6	36.8 [50]	47.7±1.4	46.4	44.2±2.3	47.9	[6]	8
ANLI R2	33.3	48.3 ^b	39.9	37.5 [5]	35.4	34.0 [50]	43.9±1.3	44.0	41.6±1.4	41.1	[6]	8
ANLI R3	33.3	43.5 ^b	39.3	40.7 [5]	34.5	40.2 [50]	47.0±1.3	48.5	42.8±2.2	46.8	[6]	8
CB	33.3	93.6 ^a	42.9	34.4 [5]	46.4	82.1 [32]	64.1±14.7	83.9	82.6±4.4	82.1	[7]	10
MNLI-m	33.3	92.2 ^a	35.7	43.7 [5]	—	—	51.1±6.2	61.2	60.8±3.7	63.5	[10]	10
MNLI-mm	33.3	91.9 ^a	37.0	43.8 [5]	—	—	51.0±6.5	62.4	61.0±3.5	63.5	[10]	10
QNLI	50.0	96.9 ^a	50.6	55.7 [5]	—	—	59.6±4.9	66.4	62.0±1.7	63.3	[12]	9
RTE	50.0	92.5 ^a	73.3	70.8 [5]	63.5	72.9 [32]	78.3±7.9	84.1	79.9±6.9	84.5	[8]	10
SNLI	33.3	91.3 ^b	33.3	54.7 [5]	—	—	43.0±7.4	53.4	62.3±2.4	65.6	[15]	9
WNLI	50.0	94.5 ^a	56.3	64.8 [5]	—	—	61.0±10.6	74.6	55.4±11.0	70.4	[14]	10
READING COMP.												
BoolQ	50.0	91.2 ^a	81.0	80.0 [1]	60.5	77.5 [32]	80.2±3.1	82.9	83.6±0.8	84.6	[4]	9
DROP (F1)	—	80.5 ^b	3.8	10.3 [1]	23.6 [†]	36.5 [20]	21.9±0.9	22.7	22.3±1.1	23.9	[2]	7
MultiRC (F1)	—	88.1 ^a	60.0	59.6 [5]	72.9	74.8 [32]	74.5±3.7	77.5	69.2±3.2	72.1	[1]	8
OBQA	25.0	85.4 ^a	41.8	50.6 [10]	57.6	65.4 [100]	77.4±1.3	78.4	77.2±1.3	78.2	[16]	7
SQuADv1 (F1)	—	96.2 ^a	22.7	50.2 [3]	—	—	79.5±1.6	80.1	82.1±0.5	82.7	[4]	8
SQuADv2 (F1)	—	83.4 ^b	11.1	34.9 [3]	59.5 [†]	69.8 [16]	40.9±1.8	44.2	40.8±0.9	43.1	[3]	10
CLOSED-BOOK QA												
ARC-challenge	25.0	81.1 ^a	42.0	49.4 [10]	51.4	51.5 [50]	61.7±1.4	63.1	63.7±0.6	63.8	[13]	7
ARC-easy	25.0	92.6 ^a	76.4	80.9 [10]	68.8	70.1 [50]	79.5±0.8	79.6	80.5±0.5	80.7	[14]	7
NQ	—	36.6 ^a	3.2	22.1 [5]	14.6	29.9 [64]	18.6±2.7	20.7	27.2±0.5	27.6	[16]	10
TQA (wiki)	—	60.5 ^a	21.9	63.3 [10]	64.3	71.2 [64]	66.5±2.6	68.1	66.5±1.0	67.3	[16]	10
TQA (tfds-dev)	—	51.0 ^a	18.4	55.1 [10]	—	—	55.0±2.3	56.7	57.2±0.6	57.8	[16]	10
COMMONSENSE												
COPA	50.0	94.8 ^a	90.0	89.0 [10]	91.0	92.0 [32]	90.6±2.0	91.0	88.5±3.8	87.0	[16]	8
HellaSwag	25.0	47.3 ^b	57.0	58.8 [10]	78.9	79.3 [20]	56.4±0.5	56.7	59.4±0.2	59.2	[3]	8
PIQA	50.0	66.8 ^b	80.3*	80.2* [10]	81.0	82.3 [50]	80.9*±0.8	80.5*	82.1*±0.3	81.7*	[10]	8
StoryCloze	50.0	89.2 ^b	79.5	83.7 [10]	83.2	87.7 [70]	92.2±1.3	93.4	93.3±0.9	94.7	[10]	8
SENTIMENT												
IMDB	50.0	95.5 ^b	76.9	83.3 [1]	—	—	94.1±0.4	94.3	94.8±0.3	95.0	[2]	7
Sentiment140	50.0	87.0 ^b	41.4	63.3 [5]	—	—	69.9±2.5	73.5	68.7±1.2	69.3	[16]	6
SST-2	50.0	97.5 ^a	51.0	92.3 [5]	71.6	95.6 [8]	92.6±1.7	94.6	94.4±0.8	94.6	[16]	8
Yelp	50.0	98.1 ^b	84.7	89.6 [3]	—	—	97.8±0.2	98.1	97.9±0.1	98.0	[4]	7
PARAPHRASE												
MRPC	50.0	90.4 ^a	53.7	64.0 [5]	—	—	69.1±1.3	69.1	67.5±1.7	67.2	[10]	10
QQP	50.0	90.6 ^a	34.9	58.9 [3]	—	—	72.1±6.8	75.9	73.5±2.9	75.9	[16]	7
PAWS Wiki	50.0	91.9 ^a	45.5	53.5 [5]	—	—	61.5±6.5	69.4	66.2±0.9	70.2	[10]	10
COREFERENCE												
DPR	50.0	84.8 ^b	54.6	57.3 [5]	—	—	60.3±3.5	66.8	62.4±1.6	63.3	[16]	10
Winogrande	50.0	65.8 ^b	68.3	68.4 [10]	70.2	77.7 [50]	67.3±2.5	71.2	72.3±0.9	72.8	[16]	10
WSC273	50.0	70.0 ^b	81.0	61.5 [5]	88.3	88.5 [32]	80.8±3.7	—	— ± —	—	[—]	10
READ. COMP. W/ COMMONSENSE												
CosmosQA	25.0	67.1 ^b	34.1	33.8 [5]	—	—	58.4±1.3	60.6	56.7±1.3	56.0	[5]	8
ReCoRD	—	93.4 ^a	87.8*	87.6* [1]	90.2	89.0 [32]	67.8*±3.0	72.5*	77.0*±2.0	79.0*	[1]	10

Table 5: Results for eight NLU task clusters. All values shown are for accuracy (or exact match) unless otherwise indicated. [k] indicates the number of few-shot exemplars. #t indicates the number of templates that FLAN is evaluated on. ^aT5-11B, ^bBERT-large. *see data contamination (Appendix D). WSC273 (Levesque et al., 2012) does not have training or validation sets, and so we do not compute few-shot results for FLAN. For Trivia QA (TQA), we report exact match (EM) on both the wikipedia subset of the dev set to compare with GPT-3, as well as the full TFDS dev set.

		FLAN 137B											
		Supervised Model	Base LM 137B		GPT-3 175B		zero-shot		few-shot		[k]	#t	
Metric			zero-shot	few-shot [k]	zero-shot	few-shot [k]	average template	best dev template	average template	best dev template			
TRANSLATION													
WMT '14 En→Fr	BLEU	35.0 ^d	11.2	31.5 [5]	25.2	32.6 [64]	32.9±1.1	33.9	33.9±0.2	33.8	[9]	5	
WMT '14 Fr→En	BLEU	45.6 ^c	7.2	34.7 [5]	21.2	39.2 [64]	35.5±1.3	35.9	38.0±0.1	37.9	[9]	3	
WMT '16 En→De	BLEU	38.6 ^f	7.7	26.7 [5]	24.6	29.7 [64]	25.4±1.8	27.0	26.8±0.4	26.1	[11]	5	
WMT '16 De→En	BLEU	41.2 ^e	20.8	36.8 [5]	27.2	40.6 [64]	38.9±0.3	38.9	40.6±0.1	40.7	[11]	3	
WMT '16 En→Ro	BLEU	39.9 ^g	3.5	22.9 [5]	14.1	21.0 [64]	16.7±1.6	18.9	20.5±0.1	20.5	[9]	5	
WMT '16 Ro→En	BLEU	38.5 ^g	9.7	37.5 [5]	19.9	39.5 [64]	36.8±0.5	37.3	38.2±0.1	38.1	[9]	3	
STRUCT TO TEXT													
CommonGen	Rouge-1	64.0 ^a	3.9	56.7 [3]	—	—	54.6±2.3	56.3	56.6±0.3	56.4	[16]	6	
	Rouge-2	29.4 ^a	1.5	29.6 [3]	—	—	28.8±2.4	27.6	30.9±0.7	29.9	[16]	6	
	Rouge-L	54.5 ^a	3.2	48.5 [3]	—	—	48.4±1.9	48.7	50.7±0.2	51.0	[16]	6	
DART	Rouge-1	—	11.3	56.0 [3]	—	—	45.5±4.2	48.9	57.9±1.6	59.2	[11]	7	
	Rouge-2	—	1.5	29.6 [3]	—	—	25.0±3.7	30.0	35.8±1.0	36.2	[11]	7	
	Rouge-L	—	3.2	48.5 [3]	—	—	38.4±3.8	43.4	48.5±0.9	48.2	[11]	7	
E2ENLG	Rouge-1	72.6 ^a	6.2	56.7 [3]	—	—	44.8±3.9	51.4	59.1±1.3	59.7	[12]	9	
	Rouge-2	47.5 ^a	2.5	31.4 [3]	—	—	24.2±3.6	30.1	33.2±1.1	33.6	[12]	9	
	Rouge-L	56.4 ^a	4.9	41.1 [3]	—	—	37.0±3.5	42.4	44.9±0.8	45.1	[12]	9	
WebNLG	Rouge-1	83.5 ^a	13.9	68.3 [3]	—	—	50.6±4.7	57.7	68.5±2.2	71.2	[10]	8	
	Rouge-2	63.6 ^a	6.9	46.0 [3]	—	—	29.8±4.2	35.4	48.0±1.5	49.8	[10]	8	
	Rouge-L	71.0 ^a	11.8	56.5 [3]	—	—	43.4±4.5	49.7	58.8±1.1	60.2	[10]	8	

Table 6: Results for translation and struct-to-text tasks. [k] indicates the number of few-shot exemplars. #t indicates the number of templates that FLAN is evaluated on. ^aT5-11B, ^cEdunov et al. (2018), ^dDurrani et al. (2014), ^eWang et al. (2019b), ^fSennrich et al. (2016), ^gLiu et al. (2020).

For struct to text, reported T5/mT5 results are from the GEM benchmark paper (Gehrmann et al., 2021), though we do not report their results for DART (through correspondence with authors, we confirmed that their results for DART were incorrect). Though we use a summarization task cluster during instruction tuning, we leave evaluation of summarization for future work, as the mean input of most summarization datasets exceeds FLAN’s input length of 1024 tokens.

B FURTHER ABLATION STUDY: DATASETS PER TASK CLUSTER & TEMPLATES PER DATASET

Our primary hypothesis is that instruction tuning on a diverse set of tasks improves performance on unseen tasks. §4.1 showed that adding more task clusters improves performance; here, we further explore whether adding additional datasets improves performance when the number of task clusters is held constant. We use the same split as in §4.1, where the NLI, commonsense reasoning, and closed-book QA clusters are held-out, and seven other task clusters remain for instruction tuning. For these seven task clusters, we instruction tune models using just one dataset per task cluster and using four datasets per task cluster (for task clusters that did not have four tasks, we just used all available tasks). In addition, we simultaneously explore the role of the number of instruction templates per dataset; as mentioned in §2.1, for each dataset we manually composed ten instructional templates for instruction tuning. Here, we instruction tune models using 1, 4, and 10 templates per dataset.

Figure 8 shows these results. Using more datasets per cluster improved performance by almost 10% on average across the three held-out clusters. Using more templates per dataset, however, had a comparatively negligible effect on performance when there was one task per cluster, which disappeared when there were four tasks per cluster. The small effect of templates is striking given our original motivation that composing ten templates per task would mitigate overfitting to any particular template. This results serves to underscore, however, the unpredictability of finetuning large language models, as one hypothesis is that models at such scale do not easily overfit to a finetuning single task.




Figure 8: Effect of datasets per task cluster and templates per dataset on performance on three held-out clusters: NLI, commonsense reasoning, and closed-book QA. Adding more datasets per task cluster substantially improves performance. Using more templates per dataset, however, only had a very small effect on performance, which disappeared when there were sufficient dataset per task cluster.

C FURTHER ANALYSIS: INSTRUCTION TUNING FACILITATES PROMPT TUNING

As we’ve seen that instruction tuning improves the ability of a model to respond to instructions, it follows that, if FLAN is indeed more amenable to performing NLP tasks, then it should also achieve better performance when performing inference using soft prompts, represented by prepended continuous variables optimized via prompt tuning (Li & Liang, 2021; Lester et al., 2021). As further analysis, we train continuous prompts for each of the SuperGLUE (Wang et al., 2019a) tasks in accordance with the cluster splits from §2.2 such that when prompt-tuning on task \mathcal{T} , no tasks in the same cluster as \mathcal{T} were seen during instruction tuning. Our prompt tuning setup follows the procedure of Lester et al. (2021) except that we use a prompt length of 10, weight decay of $1e-4$, and did not use dropout on the attention scores; we found in preliminary experiments that these changes improved the performance of Base LM.

Table 7 shows the results of these prompt tuning experiments for both using a fully-supervised training set and in a low-resource setting with only 32 training examples. We see that in all scenarios, prompt tuning works better with FLAN than Base LM. In many cases, especially for the low-resource setting, prompt tuning on FLAN even achieves more than 10% improvement over prompt tuning on the Base LM. This result exemplifies in another way how instruction tuning can result in a checkpoint that is more desirable for performing NLP tasks.

PROMPT TUNING ANALYSIS									
	Prompt tuning train. examples	BoolQ	CB	CoPA	MultiRC	ReCoRD	RTE	WiC	WSC
		acc.	acc.	acc.	F1	acc.	acc.	acc.	acc.
Base LM	32	55.5	55.4	87.0	65.4	78.0	52.4	51.6	65.4
FLAN		77.5	87.5	91.0	76.8	80.8	83.0	57.8	70.2
Base LM	full dataset	82.8	87.5	90.0	78.6	84.8	82.0	54.9	72.7
FLAN		86.3	98.2	94.0	83.4	85.1	91.7	74.0	86.5

Table 7: FLAN responds better to continuous inputs attained via prompt tuning than Base LM. When prompt tuning on a given dataset, no tasks from the same cluster as that dataset were seen during instruction tuning.

D DATA CONTAMINATION ANALYSIS

One reasonable concern is that since the pretraining corpus of FLAN has more than 2 trillion tokens, it is possible that examples from a given evaluation dataset may have already been seen verbatim by the model during pre-training, hence inflating the performance of our purported zero-shot model. To this end, like GPT-3 (Brown et al., 2020), we perform post-hoc data contamination analysis to investigate whether the performance of the model is in fact inflated by evaluating on examples that occurred in the pretraining dataset.

Our data contamination procedure follows the setup of Brown et al. (2020), which, for each benchmark, produces a “clean” version that removes all potentially leaked examples, defined as examples for which any n -gram (n varies per dataset but is roughly 13) overlapped with anything in the pretraining corpus. We use the same n per dataset as Brown et al. (2020), only differing in that we use the SentencePiece tokenizer (Kudo & Richardson, 2018) instead of splitting on spaces. We then evaluate our model on this clean subset, comparing against model performance on the original dataset (clean + dirty). Lower performance on the clean subset would suggest that data contamination leads to inflated results.

Figure 9 summarizes these results, with the exact numbers given in Table 8. We see several trends very similar to those in the GPT-3 paper: (1) many datasets had a substantial number of examples that overlapped with the pretraining data, (2) across all datasets, we do not see a correlation that evaluating on clean data does worse than evaluating on the total dataset, and (3) as datasets had fewer clean examples, there was higher variance in the percent change in performance (likely due to a smaller number of clean examples).




Figure 9: Like GPT-3, we also measured performance on cleaned versions of our datasets, which had high confidence to be unseen in the pretraining data of FLAN. We do not see a correlation that FLAN performed better on evaluation sets for which examples occurred more often in the pretraining data. When the percent of clean data is very small, there are fewer examples for computing the clean performance, which leads to high variance.

Like GPT-3, we also found that DROP and SQuADv2 had almost total overlap with the pretraining data. We follow their procedure of manually inspecting the data, and find that most overlapping n -grams were only in the contexts of examples (99.6% for DROP and 97.2% for SQuADv2). Overlaps never occurred in both the question and answer for DROP, and only occurred for both the question and answer for SQuADv2 in 5 of the 11,153 evaluation examples. Hence, for these two datasets, the model gains only background information and cannot memorize the answer to any specific questions (aside from the five examples in SQuADv2).

ANLI R1 and R2 (Nie et al., 2020) also had almost complete data contamination, to a much higher degree than GPT-3. Upon further inspection, we see that most overlaps occur in example contexts and not hypotheses (97.3% for ANLI R1 and 98.2% for ANLI R2). As ANLI R1 and R2 are based entirely from Wikipedia examples (R3 is not), we posit that this higher degree of contamination in our pretraining dataset compared with GPT-3’s is potentially due to using a more-recent version of Wikipedia that includes the contexts used in ANLI R1 and R2 (which were collected in 2019).

Because seeing a particular context in pretraining does not help with the NLI task given a new, unseen sentence, we think it is unlikely that these overlaps affected performance on the two datasets.

Of the remaining datasets, only ReCoRD and PIQA had a clean subset performance that was lower than the overall evaluation set performance by more than 1%. These two datasets are language modeling (i.e., “what’s the best continuation of this sentence?”), and so it is more likely compared with previous tasks that seeing a complete sentence in the pretraining data could help the model predict the right answer in downstream evaluations. For PIQA, both the goal and solution had overlaps in 93 of the 1,838 evaluation examples, and for ReCoRD, the query had overlaps in 2,320 of the 10,000 training examples. We hence mark these results with an asterisk * in Table 4 and Table 5. Brown et al. (2020) also reported substantial contamination rates for these two datasets (61% dirty for ReCoRD and 29% for PIQA), and also mark PIQA with an asterisk.

As this overlap analysis follows that performed in Brown et al. (2020), we reiterate the same caveats: the conservative nature of our n -gram matching procedure likely introduces additional false positives; there are no guarantees that the clean subset is drawn from the same distribution as the overall subset; and, accurately detecting test contamination is a relatively new research area without established best practices. Moreover, as our pretraining corpus is almost five times larger than that used for GPT-3 (which was 500B tokens), it is possible that there are more false positives in detecting dirty data.

Dataset	Metric	Total count	Total acc/F1/BLEU	Clean count	Clean acc/F1/BLEU	% clean	% Diff (clean – overall)
DROP	F1	9,536	22.4	61	33.0	0.6	47.4
SQuADv2	F1	11,873	41.3	106	38.7	0.9	-6.2
ANLI R1	acc	1,000	48.1	14	57.1	1.4	18.8
ANLI R2	acc	1,000	42.9	21	38.1	2.1	-11.2
ReCoRD	acc	10,000	4.6	3,203	4.5	32.0	-2.7
MultiRC	acc	4,848	75.4	1,972	75.7	40.7	0.5
PIQA	acc	1,838	23.7	896	23.3	48.7	-1.7
ANLI R3	acc	1,200	44.2	718	45.3	59.8	2.5
HellaSwag	acc	10,042	28.5	6,578	28.7	65.5	0.7
RTE	acc	2,77	84.1	183	84.2	66.1	0.0
WMT’14 En→Fr	BLEU	3,003	31.3	2,243	31.5	74.7	0.9
WMT’14 Fr→En	BLEU	3,003	34.0	2,243	34.1	74.7	0.2
BoolQ	acc	3,270	76.5	2,515	76.3	76.9	-0.4
TQA (tfds-dev)	F1	11,313	62.2	8,731	62.0	77.2	-0.2
ARC Easy	acc	2,365	79.5	1,888	79.0	79.8	-0.6
ARC Challenge	acc	1,165	63.1	983	64.2	84.4	1.7
OpenbookQA	acc	500	74.6	425	74.8	85.0	0.3
WMT’16 En→De	BLEU	2,999	22.7	2,569	23.0	85.7	1.4
WMT’16 De→En	BLEU	2,999	38.6	2,569	38.7	85.7	0.2
WMT’16 En→Ro	BLEU	1,999	15.5	1,752	15.4	87.6	-0.7
WMT’16 Ro→En	BLEU	1,999	36.7	1,752	36.8	87.6	0.1
COPA	acc	100	88.0	91	87.9	91.0	-0.1
CB	acc	56	41.1	53	41.5	94.6	1.1
NQ	F1	3,610	24.5	3,495	24.3	96.8	-0.5
StoryCloze	acc	1,871	92.1	1,864	92.1	99.6	0.0
Winogrande	acc	1,267	39.4	1,265	39.4	99.8	0.2

Table 8: Overlap statistics for the subset of datasets that are also used in GPT-3, sorted from dirtiest to cleanest. An evaluation example was dirty if it had any n -gram collision with the pretraining corpus. We computed these results for FLAN’s performance using only a single template for each dataset, so they differ slightly compared with the results for average performance over all templates.

E EXTENDED RELATED WORK

E.1 LANGUAGE MODELS AND MULTI-TASK LEARNING

Our work is broadly inspired by a long line of prior work on language models for NLP applications (Dai & Le, 2015; Peters et al., 2018; Howard & Ruder, 2018; Radford et al., 2018; 2019, *inter alia*). Instruction tuning can be seen as a formulation of multitask learning (MTL), which is an established area within deep learning (Collobert et al., 2011; Luong et al., 2016; Ruder, 2017; Velay & Daniel, 2018; Clark et al., 2019b; Liu et al., 2019b, *inter alia*)—see Worsham & Kalita (2020) for a recent survey on MTL for NLP. Differing from prior MTL work which focuses on performance improvements across training tasks (Raffel et al., 2020; Aghajanyan et al., 2021) or to new domains (Axelrod et al., 2011), our work is motivated by improving zero-shot generalization to tasks that were not seen in training.

E.2 ZERO-SHOT LEARNING AND META-LEARNING

Our work also falls in the well-established category of zero-shot learning, which has historically been used to refer to classifying instances among a set of unseen categories (Lampert et al., 2009; Romera-Paredes & Torr, 2015; Srivastava et al., 2018; Yin et al., 2019, *inter alia*). In NLP, zero-shot learning work also includes translating between unseen language pairs (Johnson et al., 2017; Pham et al., 2019), language modeling on unseen languages (Lauscher et al., 2020), as well as various NLP applications (Liu et al., 2019a; Corazza et al., 2020; Wang et al., 2021; Reif et al., 2021). Most recently, the emergent ability of language models (Brown et al., 2020) has led to increased interest in how models generalize to unseen tasks, the definition of zero-shot learning used in our paper. In addition, meta-learning (Finn et al., 2017; Vanschoren, 2018, *inter alia*) also broadly tries to train models that adapt quickly to unseen tasks, typically based on a few examples.

E.3 PROMPTING

Instruction tuning leverages the intuition that language models at scale contain substantial world knowledge and can perform a range of NLP tasks (Brown et al., 2020, see also Bommasani et al. (2021)). Another line of work that shares this goal prompts models with continuous inputs optimized via backpropagation to substantially improve performance (Li & Liang, 2021; Lester et al., 2021; Qin & Eisner, 2021). Although the success of these approaches depends heavily on model scale (Lester et al., 2021), for which large models can be costly to serve, the ability of a single large model to perform many tasks slightly eases this burden. As shown by our experiments in Appendix C, prompt tuning is an orthogonal method for which instruction tuning can additionally improve performance.

Our work shares similar motivations with prompting in that we use inference-time text interactions to prompt a single model, without creating separate checkpoints for each task. Whereas prompting work such as GPT-3 uses prompt engineering to write prompts that intentionally mimic text that is likely to be seen during pretraining (e.g., for MultiRC GPT-3 tries a prompt that mimics a test with an answer key), we hope that finetuning models to respond to natural language instructions instead of completing a sentence will make such large models more accessible to non-technical users.

E.4 FINETUNING LARGE LANGUAGE MODELS

Finetuning pretrained language models is a well-established method in NLP, with much of the work so far occurring on models in the range of 100M to 10B parameters (Dai & Le, 2015; Devlin et al., 2019; Raffel et al., 2020; Lewis et al., 2020, *inter alia*). For models of $O(100B)$ parameters, recent work has finetuned task-specific models for program synthesis (Austin et al., 2021; Chen et al., 2021), summarization (Wu et al., 2021), as well as improved bias and fairness behavior (Solaiman & Dennison, 2021). In addition to the traditional “dense” models, sparse mixture of experts (MoE) models of up to more than 1T parameters have been trained and finetuned (Lepikhin et al., 2020; Fedus et al., 2021). Compared with this prior work that finetunes and evaluates on the same downstream task, our setup studies the effect of instruction tuning on ability to perform unseen tasks.

E.5 MULTI-TASK QUESTION ANSWERING

The instructions we use for instruction tuning are similar to QA-based task formulation research, which aims to unify NLP tasks by casting them as question-answering over a context. For instance, McCann et al. (2018) cast ten NLP tasks as QA and train a model on a collection of tasks formulated with natural language prompts; they report transfer learning gains on finetuning tasks as well as zero-shot domain adaptation results on SNLI (Bowman et al., 2015) and Amazon/Yelp Reviews (Kotzias et al., 2015). While McCann et al. (2018) does not leverage unsupervised pre-training and only reports zero-shot transfer to unseen domains, our work uses a pretrained LM and focuses on zero-shot performance on unseen task clusters. Focusing on binary text classification, Zhong et al. (2021) finetune T5-770M on 43 tasks phrased as yes/no questions and study the zero-shot performance on unseen tasks. In comparison, our paper is much larger in scope, empirically demonstrating the idea on a wide range of tasks with a much larger model. Other work has used QA-based task formulation for more-targeted applications including semantic role labeling (He et al., 2015), relation extraction (Levy et al., 2017), coreference resolution (Wu et al., 2020) and named entity recognition (Li et al., 2020) as question answering.

E.6 INSTRUCTIONS-BASED NLP

Recent improvements in the capabilities of language models have led to increased interest in a nascent area of instructions-based NLP (Goldwasser & Roth, 2014, and see McCarthy (1960)). Schick & Schütze (2021) (also see Gao et al., 2021; Tam et al., 2021) use task descriptions in cloze-style phrases to help language models assign soft labels for few-shot and semi-supervised learning, though this line of work finetunes new checkpoints for each downstream task. Efrat & Levy (2020) evaluated GPT-2 (Radford et al., 2019) on simple tasks ranging from retrieving the n th word of a sentence to generating examples for SQuAD, concluding that GPT-2 performs poorly across all tasks.

In terms of the setup of finetuning on a large number of tasks and evaluating on unseen tasks, two recent papers are very similar to ours. Mishra et al. (2021) finetune BART (Lewis et al., 2020) using instructions and few-shot examples for tasks such as question answering, text classification, and text modification, and find that this few-shot finetuning with instructions improves performance on unseen tasks. Ye et al. (2021) introduce a setup for cross-task few-shot learning, finding that multi-task meta-learning using MAML (Finn et al., 2017) improves the few-shot capabilities of BART on unseen downstream tasks. Our work differs from these two papers in that we focus on zero-shot learning, for which we observe the crucial importance of model scale (FLAN is 1,000x larger than BART-base).

Finally, OpenAI has a beta product called the Instruct Series. Though they have not released a research paper on it at this time, based on examples posted on their website,⁴ their product appears to have goals similar to those in our work.

⁴<https://beta.openai.com/docs/engines/instruct-series-beta>

CHANGES TO V2 (THIS VERSION) FROM V1

These changes include the following:

- We updated the terminology to “datasets” and “task clusters.”
- We renamed the previous “open-domain QA” task cluster to “closed-book QA.”
- We extended the related work section and moved it to the Appendix E, using a shorter version in the main body.
- We added FLAN and Base LM results for additional datasets for which GPT-3 results were not reported. These are in Appendix A.2.
- For TriviaQA, v1 reported results on the tfds dev set of 11,313 examples. GPT-3 actually evaluates on the wikipedia dev set of 7,993 examples, so we ran an additional evaluation on that dev set in order to compare with GPT-3’s performance. Zero-shot FLAN now beats zero-shot GPT-3 on that task (and therefore on 20 of 25 tasks). We still show the original result in Table 5, though there is no GPT-3 result to compare with.
- We moved commonsense reasoning and coreference resolution from the main body to Appendix A.1.
- We moved prompt tuning from the main body to Appendix C.
- We added data contamination analysis (Appendix D).
- We added few-shot instruction tuning (§4.3).
- We cited additional datasets in Appendix F.

F TASKS AND DATASETS

This appendix further details the datasets that we use in this paper. We group datasets into one of the following task clusters:

- **Natural language inference** concerns how two sentences relate, typically asking, given a first sentence, whether a second sentence is true, false, or possibly true. We use the following datasets:
 1. ANLI (Nie et al., 2020)
 2. CB (De Marneffe et al., 2019)
 3. MNLI (Williams et al., 2018)
 4. QNLI (Rajpurkar et al., 2018)
 5. SNLI (Bowman et al., 2015)
 6. WNLI (Levesque et al., 2012)
 7. RTE (Dagan et al., 2005; Haim et al., 2006; Giampiccolo et al., 2007; Bentivogli et al., 2009)
- **Reading comprehension** tests the ability to answer a question when given a passage that contains the answer. We use the following datasets:
 1. BoolQ Clark et al. (2019a)
 2. DROP (Dua et al., 2019)
 3. MultiRC (Khashabi et al., 2018)
 4. OBQA (Mihaylov et al., 2018)
 5. SQuADv1 (Rajpurkar et al., 2016)
 6. SQuADv2 (Rajpurkar et al., 2018)
- **Commonsense reasoning** evaluates the ability to perform physical or scientific reasoning with an element of common sense. We use the following datasets:
 1. COPA (Roemmele et al., 2011)
 2. HellaSwag (Zellers et al., 2019)
 3. PiQA (Bisk et al., 2020)
 4. StoryCloze (Mostafazadeh et al., 2016)
- **Sentiment analysis** is a classic NLP task aims to understand whether a piece of text is positive or negative. We use the following datasets:
 1. IMDB (Maas et al., 2011)
 2. Sentiment140 (Go et al., 2009)
 3. SST-2 (Socher et al., 2013)
 4. Yelp (Fast.AI)
- **Closed-book QA** asks models to answer questions about the world without specific access to information that contains the answer. We use the following datasets:
 1. ARC (Clark et al., 2018)
 2. NQ (Lee et al., 2019; Kwiatkowski et al., 2019)
 3. TriviaQA Joshi et al. (2017)
- **Paraphrase detection** asks a model to determine whether two sentences are semantically equivalent.⁵ We use the following datasets:
 1. MRPC (Dolan & Brockett, 2005)
 2. QQP (Wang et al., 2018, see)
 3. Paws Wiki (Zhang et al., 2019)
- **Coreference resolution** tests the ability to identify expressions of the same entity in some given text. We use the following datasets:
 1. DPR (Rahman & Ng, 2012)
 2. Winogrande (Sakaguchi et al., 2020)

⁵Although paraphrasing can be seen as positive entailment in both directions, it has been distinct from NLI in the academic literature.

3. WSC273 (Levesque et al., 2012)
- **Reading comprehension with commonsense** combines elements of both reading comprehension with commonsense. We use the following datasets:
 1. CosmosQA (Huang et al., 2019)
 2. ReCoRD (Zhang et al., 2018)
 - **Struct to text** tests the ability to describe some structured data using natural language. We use the following datasets:
 1. CommonGen (Lin et al., 2020)
 2. DART (Nan et al., 2021)
 3. E2ENLG (Dušek et al., 2019)
 4. WebNLG (Gardent et al., 2017)
 - **Translation** is the task of translating text from one language into a different language. We use the following datasets:
 1. En-Fr from WMT’14 (Bojar et al., 2014)
 2. En-De, En-Tr, En-Cs, En-Fi, En-Ro, and En-Ru from WMT’16 (Bojar et al., 2016)
 3. En-Es from Paracrawl (Bañón et al., 2020)
 - **Summarization** asks models to read a piece of text and generate an abbreviated summary of it. We use the following datasets:
 1. AESLC (Zhang & Tetreault, 2019)
 2. CNN-DM (See et al., 2017)
 3. Gigaword (Napoles et al., 2012)
 4. MultiNews (Fabbri et al., 2019)
 5. Newsroom (Grusky et al., 2018)
 6. Samsun (Gliwa et al., 2019)
 7. XSum (Narayan et al., 2018)
 8. AG News (Zhang et al., 2015)
 9. Opinion Abstracts - Rotten Tomatoes (Wang & Ling, 2016)
 10. Opinion Abstracts - iDebate (Wang & Ling, 2016)
 11. Wiki Lingua English (Ladhak et al., 2020)
 - Additional datasets that we assign to a miscellaneous task cluster include:
 1. Conversational question-answering: QuAC (Choi et al., 2018) and CoQA (Reddy et al., 2019)
 2. Evaluating context-sentence word meanings: WiC (Pilehvar & Camacho-Collados, 2019)
 3. Question classification: TREC (Li & Roth, 2002; Hovy et al., 2001)
 4. Linguistic acceptability: CoLA (Warstadt et al., 2019)
 5. Math questions (Saxton et al., 2019)

For all tasks, our finetuning and evaluation code uses tensorflow datasets (TFDS) to load and process datasets. Regarding the number of training examples per dataset, we limited the training set size per dataset to 30,000 so that no dataset dominated the finetuning distribution. When a test set with labels was available in TFDS, we used it; otherwise, we used the TFDS validation set as our test set, splitting the training set into a train and dev set.

On the following pages, we show inputs and outputs for evaluation tasks where we compared with GPT-3. See the attached supplementary material for the templates for all other datasets.

F.1 NATURAL LANGUAGE INFERENCE

<p>INPUT</p> <p>Joey Heindle (born 14 May 1993 in Munich) is a German singer. He is best known for winning the seventh season of the game show Ich bin ein Star – Holt mich hier raus! and finishing in 5th place in season 9 of Deutschland sucht den Superstar, despite universally negative reviews from the jury each week.</p> <p>Based on the paragraph above can we conclude that "Joey Heindle was highly disliked by people on television."?</p> <p>OPTIONS:</p> <ul style="list-style-type: none"> - Yes - It's impossible to say - No
<p>TARGET</p> <p>Yes</p>

Table 9: Example input and target for Adversarial NLI (ANLI). ANLI (Nie et al., 2020) is a large-scale NLI benchmark with adversarial examples collected iteratively with a human and model in the loop. The task is to determine whether a hypothesis is entailed by a premise (entailment, not entailment, or impossible to say). There are three rounds, R1–R3. Of the three training sets with 16,946, 45,460, and 100,459 examples, we use 16,946, 30,000, and 30,000 for train and 200 from each of the three TFDS validation sets for dev. We use the TFDS “test” sets of 1,000, 1,000, and 1,200 examples as our test set for reporting numbers.

<p>INPUT</p> <p>A: so I watch the fish, you know. Whatever I can do to keep myself occupied. I like to have the TV on, because that usually keeps me, um, more occupied. It kind of takes the time away and I don't realize, that's really the only time I ever watch TV, is when I'm on the bike. and then usually after I'm done riding the bike, just to cool myself down, I usually take a walk, you know, and that just kind of uh, gets me, you know, to where I'm not quite as tired I guess. But it's definitely a task. B: You think so? A: I can't say that I really enjoy it.</p> <p>Based on the paragraph above can we conclude that "she really enjoys it"?</p> <p>OPTIONS:</p> <ul style="list-style-type: none"> - Yes - No - It's impossible to say
<p>TARGET</p> <p>No</p>

Table 10: Example input and target for Commitment Bank (CB). CB (De Marneffe et al., 2019) is a corpus of texts in which a hypothesis is extracted from a premise, and the task is to determine whether the hypothesis is entailed by the premise (entailment, not entailment, or impossible to say). Of the training set with 250 examples, we use 200 for train and 50 for dev. We use the TFDS validation set of 56 examples as our test set for reporting numbers.

<p><u>INPUT</u></p> <p>After years of study, the Vatican's doctrinal congregation has sent church leaders a confidential document concluding that "sex-change" procedures do not change a person's gender in the eyes of the church.</p> <p>Based on the paragraph above can we conclude that "Sex-change operations become more common."?</p> <p>OPTIONS:</p> <ul style="list-style-type: none"> - yes - no
<p><u>TARGET</u></p> <p>no</p>

Table 11: Example input and target for Recognizing Textual Entailment (RTE). RTE (Dagan et al., 2005; Haim et al., 2006; Giampiccolo et al., 2007; Bentivogli et al., 2009) asks whether a second sentence is entailed by a first (binary, either entailed or not entailed). Of the training set with 2490 examples, we use 2,290 for train and 200 for dev. We use the TFDS validation set of 277 examples as our test set for reporting numbers.

F.2 READING COMPREHENSION

<p>INPUT</p> <p>There are four ways an individual can acquire Canadian citizenship: by birth on Canadian soil; by descent (being born to a Canadian parent); by grant (naturalization); and by adoption. Among them, only citizenship by birth is granted automatically with limited exceptions, while citizenship by descent or adoption is acquired automatically if the specified conditions have been met. Citizenship by grant, on the other hand, must be approved by the Minister of Immigration, Refugees and Citizenship.</p> <p>Can we conclude that can i get canadian citizenship if my grandfather was canadian?</p> <p>OPTIONS:</p> <ul style="list-style-type: none"> - no - yes
<p>TARGET</p> <p>no</p>

Table 12: Example input and target for Boolean Questions (BoolQ). BoolQ Clark et al. (2019a) asks a yes/no question based on a passage and a question. Of the training set with 9,427 examples, we use 9,227 for train and 200 for dev. We use the TFDS validation set of 3,270 examples as our test set for reporting numbers.

<p>INPUT</p> <p>Imagine you are standing in a farm field in central Illinois. The land is so flat you can see for miles and miles. On a clear day, you might see a grain silo 20 miles away. You might think to yourself, it sure is flat around here. If you drive one hundred miles to the south, the landscape changes. In southern Illinois, there are rolling hills. Why do you think this is? What could have caused these features? There are no big rivers that may have eroded and deposited this material. The ground is capable of supporting grass and trees, so wind erosion would not explain it. To answer the question, you need to go back 12,000 years. Around 12,000 years ago, a giant ice sheet covered much of the Midwest United States. Springfield, Illinois, was covered by over a mile of ice. Its hard to imagine a mile thick sheet of ice. The massive ice sheet, called a glacier, caused the features on the land you see today. Where did glaciers go? Where can you see them today? Glaciers are masses of flowing ice.</p> <p>Question: "How big were the glaciers?"</p> <p>Response: "One mile"</p> <p>Does the response correctly answer the question?</p> <p>OPTIONS:</p> <ul style="list-style-type: none"> - no - yes
<p>TARGET</p> <p>yes</p>

Table 13: Example input and target for Multi-Sentence Reading Comprehension (MultiRC). MultiRC Khashabi et al. (2018) asks an open-ended question given a paragraph that contains the answer. Of the training set with 27,243 examples, we use 27,043 for train and 200 for dev. We use the TFDS validation set of 4,848 examples as our test set for reporting numbers.

<p><u>INPUT</u> soil is a renewable resource for growing plants A plant that needs to expand will be able to have an endless resource in</p> <p><u>OPTIONS:</u> - dirt - pesticides - pay - beans</p>
<p><u>TARGET</u> dirt</p>

Table 14: Example input and target for Openbook Question Answering (OBQA). OBQA (Mihaylov et al., 2018) asks 4-way multiple choice questions based facts. Of the training set with 4,957 examples, we use all for train and 200 in the TFDS validation set of 500 examples for dev. We use the TFDS test set of 500 examples as our test set for reporting numbers.

F.3 COMMONSENSE REASONING

<p>INPUT I packed up my belongings. What is the cause?</p> <p>OPTIONS: - I was hunting for a new apartment. - I was moving out of my apartment.</p>
<p>TARGET I was moving out of my apartment.</p>

Table 15: Example input and target for Choice of Plausible Alternatives (COPA). COPA (Roemmele et al., 2011) is a causal reasoning task that asks to infer either a cause of effect of a premise from two choices. Of the training set with 400 examples, we use 350 for train and 50 for dev. We use the TFDS validation set of 100 examples as our test set for reporting numbers.

<p>INPUT What happens next in this paragraph?</p> <p>Once the rope is inside the hook, he begins moving up the wall but shortly after he stops and begins talking. The male then begins talking about the clip again and goes back up the wall. as he</p> <p>OPTIONS: - progresses, there are hooks everywhere on the wall and when he gets near them, he puts his rope inside of it for support and safety. - changes time, an instant replay of his initial move is shown a second time. - continues to talk, another male speaks about the move and shows another closeup of the plex by the male. - continues, other people start to arrive and begin to hang out with him as he makes a few parts of the rope.</p>
<p>TARGET progresses, there are hooks everywhere on the wall and when he gets near them, he puts his rope inside of it for support and safety.</p>

Table 16: Example input and target for Commonsense Sentence Completion (HellaSwag). HellaSwag (Zellers et al., 2019) tests for sentence completion that requires common sense, asking for the most probable ending given four contexts. Of the training set with 39,905 examples, we use 30,000 for train and 200 for dev. We use the TFDS validation set of 10,042 examples as our test set for reporting numbers.

<p><u>INPUT</u> Here is a goal: Remove smell from garbage disposal.</p> <p>How would you accomplish this goal?</p> <p>OPTIONS:</p> <ul style="list-style-type: none"> - Create soda ice cubes and grind through disposal. - Create vinegar ice cubes and grind through disposal.
<p><u>TARGET</u> Create vinegar ice cubes and grind through disposal.</p>

Table 17: Example input and target for Physical Question Answering (PiQA). PiQA (Bisk et al., 2020) is a commonsense QA benchmark for naive physics reasoning, where a solution to a goal must be selected from two choices. Of the training set with 16,113 examples, we use 16,013 for train and 100 for dev. We use the TFDS validation set of 1,838 examples as our test set for reporting numbers.

<p><u>INPUT</u> Caroline never drinks carbonated beverages. Her friends pick on her because of it. One day they challenged her to drink a soda. Caroline wanted to win the challenge.</p> <p>Predict the next sentence.</p> <p>OPTIONS:</p> <ul style="list-style-type: none"> - Caroline refused to open the soda. - Caroline opened the soda and drank it all in one gulp!
<p><u>TARGET</u> Caroline opened the soda and drank it all in one gulp!</p>

Table 18: Example input and target for The Story Cloze Test (StoryCloze). StoryCloze (Mostafazadeh et al., 2016) is a commonsense reasoning framework for story generation, where a system chooses the correct ending to a four-sentence story. We use the 2016 version on TFDS. Of the validation set with 1,871 examples (no training set is available), we use 1,671 for train and 200 for dev. We use the TFDS test set of 1,871 examples as our test set for reporting numbers.

F.4 CLOSED-BOOK QA

<p><u>INPUT</u> What season is the Northern Hemisphere experiencing when it is tilted directly toward the Sun?</p> <p><u>OPTIONS:</u> - fall - winter - spring - summer</p>
<p><u>TARGET</u> summer</p>

Table 19: Example input and target for The AI2 Reasoning Challenge (ARC). ARC (Clark et al., 2018) asks grade-school level 4-way multiple choice science questions. There is a challenge set and an easy set, where the challenge set questions were answered incorrectly by both a retrieval-based algorithm and a co-occurrence algorithm. Of the training sets with 1,119 examples (challenge) and 2,251 (easy), we use we use 919 and 2,051 respectively for train and 200 each for dev. We use the TFDS test sets of 1,172 and 2,376 examples respectively as our test set for reporting numbers.

<p><u>INPUT</u> Question: who is the girl in more than you know?? Answer:</p>
<p><u>TARGET</u> Romi Van Renterghem.</p>

Table 20: Example input and target for Natural Questions (Open) (NQ). NQ (Lee et al., 2019; Kwiatkowski et al., 2019) asks for an open-ended answer given a question, where all questions can be answered using the contents of Wikipedia. Of the training set of 87,925 examples, we use 30,000 for train and 200 for dev. We use the TFDS validation set of 3,610 examples as our test set for reporting numbers.

<p><u>INPUT</u> Please answer this question: Henry Croft, an orphan street sweeper who collected money for charity, is associated with what organised charitable tradition of working class culture in London, England?</p>
<p><u>TARGET</u> pearly kings and queens</p>

Table 21: Example input and target for Trivia Question Answering (TriviaQA). TriviaQA Joshi et al. (2017) includes question-answer pairs authored by trivia enthusiasts. Of the training set of 87,622 examples, we use 30,000 for train and 200 for dev. We use 7,993 examples from Wikipedia of the 11,313 examples in the TFDS validation set, which is the same validation set used in (Brown et al., 2020). as our test set for reporting numbers.

F.5 COREFERENCE RESOLUTION

<u>INPUT</u> How does the sentence end? Elena wanted to move out of her parents fast but Victoria wanted to stay for a while, OPTIONS: - Elena went to school. - Victoria went to school.
<u>TARGET</u> Victoria went to school.

Table 22: Example input and target for Adversarial Winograd Schema Challenge (Winogrande). Winogrande (Sakaguchi et al., 2020) tests for coreference resolution by asking a model to fill in a masked token in a sentence by choosing an entity from two options. Of the 40.4k examples in the XL training set, we use 30,000 for train and 200 for dev. We use the TFDS validation set of 1,267 as our test set for reporting numbers.

<u>INPUT</u> Jane knocked on Susan’s door, but there was no answer. OPTIONS: - Jane was out. - Susan was out.
<u>TARGET</u> Susan was out.

Table 23: Example input and target for Winograd Schema Challenge (WSC273). WSC273 (Levesque et al., 2012) tests for coreference resolution by asking a model to complete the sentence in a fashion that requires understanding the entities in the sentence. Of the 0 examples in the training set (WSC273 is test-set only), we use none for train and none for dev. We use the TFDS test set as our test set for reporting numbers.

F.6 READING COMPREHENSION WITH COMMONSENSE

INPUT

Complete the passage.

(CNN) – At first glance, "The Flat" might seem like an episode of "Hoarders," Israeli-style. The documentary film opens after an elderly woman dies in Tel Aviv. Her grandchildren assemble to clean out her apartment, packed with dusty books, vintage clothing (dozens of pairs of fancy gloves, for instance), enough purses to stock a department store, jewelry, mementoes and closets full of knickknacks. But buried among the detritus they chance upon something remarkable – mysterious papers linking the grandparents to an important Nazi figure. How could such ardent Zionists, who left their native Germany in the early 1930s, have been involved with an SS official like Leopold von Mildenstein?

What I found out was this journey, the Nazi (

OPTIONS:

- Arnon Goldfinger) and his wife were accompanied by my grandparents," Goldfinger told CNN.
- CNN) and his wife were accompanied by my grandparents," Goldfinger told CNN.
- Germany) and his wife were accompanied by my grandparents," Goldfinger told CNN.
- Israeli) and his wife were accompanied by my grandparents," Goldfinger told CNN.
- Leopold von Mildenstein) and his wife were accompanied by my grandparents," Goldfinger told CNN.
- Nazi) and his wife were accompanied by my grandparents," Goldfinger told CNN.
- SS) and his wife were accompanied by my grandparents," Goldfinger told CNN.
- Tel Aviv) and his wife were accompanied by my grandparents," Goldfinger told CNN.
- The Flat) and his wife were accompanied by my grandparents," Goldfinger told CNN.
- Zionists) and his wife were accompanied by my grandparents," Goldfinger told CNN.

TARGET

Leopold von Mildenstein) and his wife were accompanied by my grandparents," Goldfinger told CNN.

Table 24: Example input and target for Reading Comprehension with Commonsense Reasoning (ReCoRD). ReCoRD (Zhang et al., 2018) asks for the answer to a cloze-style question where an entity is masked out. Of the the training set of 100,730 examples, we use 30,000 for train and 200 for dev. We use the TFDS validation set of 10,000 examples as our test set for reporting numbers.

F.7 TRANSLATION (7 LANGUAGES)

INPUT

Here the largest town of the district is located: Nordenham , lying opposite to Bremerhaven at the Weser mouth.

Translate to German

TARGET

An der B 211 befindet sich in Loyermoor der so genannte "Geest-Abbruch", der eine Höhendifferenz von gut 30 Meter überbrückt.

Table 25: Example input and output for translation. This example is from WMT'16 English–German; all languages use the same translation templates.